



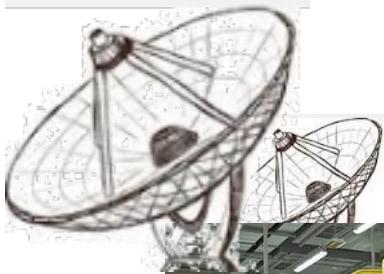
CorrelX: A Cloud-Based VLBI Correlator

V. Pankratius, A. J. Vazquez, P. Elosegui

Massachusetts Institute of Technology
Haystack Observatory

Slides based on 5th International VLBI Technology Workshop presentation
12-14 October 2016, MIT Haystack

Trends: Software-Based Instruments & Cloud Computing



Cloud Computing

Algorithms

Signal Processing

Imaging

Search & Classification

Simulations

Data analysis, mining, retrieval

Scalability is Key

- Examples: Current & Future Radio Arrays



Operational



Under construction



Planned

- 10s of Tbit/sec during cross-correlation
- 1 terabyte/day typical output data rate

Raw Data		Cost	Sensors
~72 Tbit/s	ASKAP (under construction)	\$160M	96-pixel phased array feeds on 36 12m dishes
~4 Tbit/s <small>(10 PB science-ready data / year)</small>	MeerKAT (under construction)	\$100M	64 13.5m antennas, wideband single-feed pixels
~10 Tbit/s	SKA Phase 1 (construction starting 2017)	\$900M	low-freq. component: $262144=2^{18}$ dipoles
~2.5 Pbit/s	SKA Phase 2 (construction ?)	?	3-4 million antennas

CorrelX

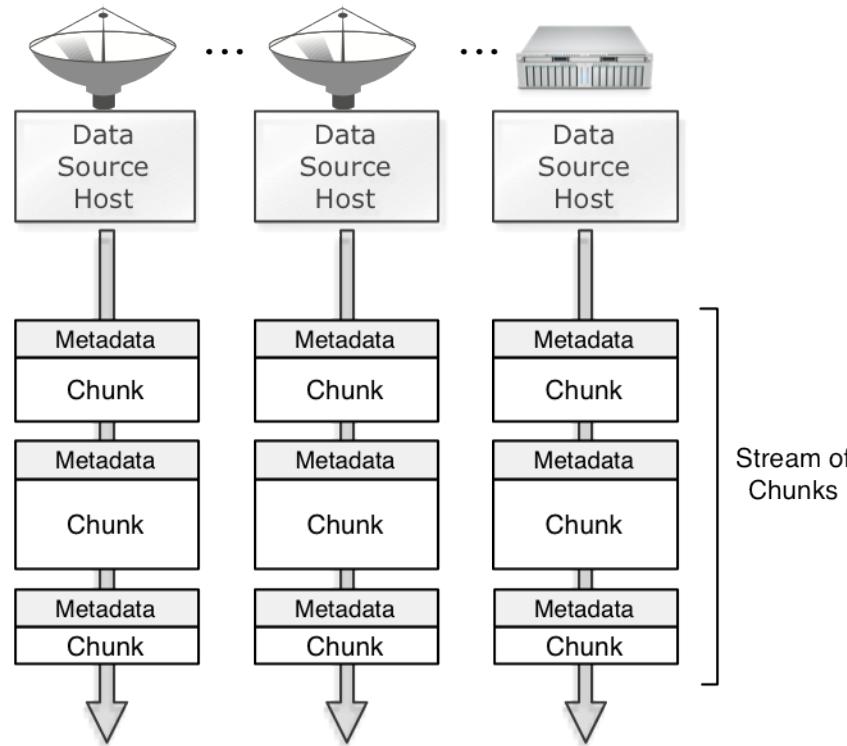
New architecture

- Breaks with current paradigms that aim to mimic hardware correlators in software
- Based on chunk streaming model
- Enables offline, out-of-order correlation of chunks

Key Features

- Cloud scalability
- Elastic correlator deployments for variable workloads, burst capability
- Reliability on large number of nodes through built-in online testing mechanisms (injected control chunks)
- Plugin-based. Enables quick prototyping and extensions for research and teaching
- Separate concerns: core correlation code vs. parallel computing code
- Leverages open-source: Map-Reduce, parallel file systems
- Enables machine-learning-based performance optimization

CorrelX: Chunk Streaming Model



- Chunks introduce locality → out-of-order parallel processing
- Enables tunable chunks (sizes, sampling, numerical precision, etc.)
- Insert control chunks with known results for validity checks on heterogeneous or unknown hardware

CorrelX

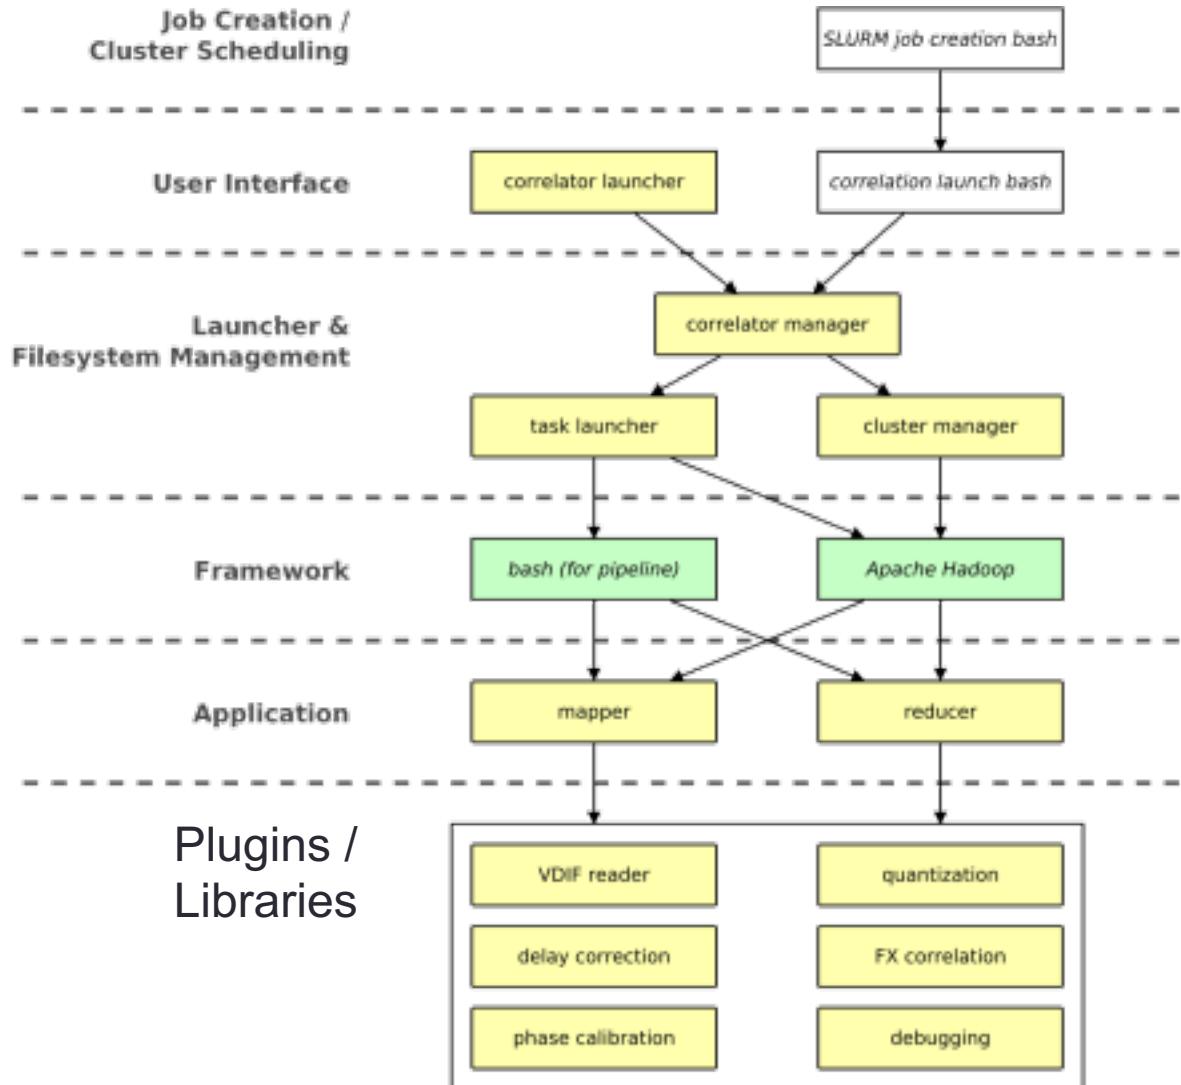
- Proof of concept prototype, incremental feature additions
- Core: **Map-Reduce paradigm**
 - Popularized by Google
 - Programming model: two types of workers:
 - Map(input_block) → set of [key,value] pairs
 - Reduce(set of [key,value] pairs) → set of [results]
 - Parallelism hidden by framework:
 - Framework launches one map per input block, one reducer per partition (i.e., all pairs identified by same key).
 - Apache Hadoop: mapper and reducer can be written in different languages (e.g., Python, C++, ...)
 - Parallel file system: Lustre
 - Max. volume size: “over 16 EB (theoretical)” (2.8.59, Oct 2016)
 - Used in over 50% of Top 100 supercomputers
 - “Presents all clients with a unified namespace for all of the files and data in the filesystem”

CorrelX

- CorrelX Correlation Configuration
 - Deployment and performance configuration: .xml files
 - Partitioner: define sub-keys for partitioning (what goes into each reducer) and sorting (to reconstruct “order” after “out-of-order” parallel computing)
- FX correlation
 - Parallelization in frequency bands and accumulation periods
 - Parallelization in baselines configurable: single-baseline-per-task, linear-scaling-stations and all-baselines-per-task (default is all-baselines-per-task)
- Delay correction using polynomials from CALC (vex2difx+calcif2).
- Phase calibration tone extraction
- Zoom bands (currently during post-processing)
- Interfaces:
 - Input format: VDIF (configuration read from frame header), complex or real, 1 or 2 bits per sample, multiple threads xor multiple channels per frame
 - Output format: text file with visibilities and phase calibration results

CorrelX Implementation Details

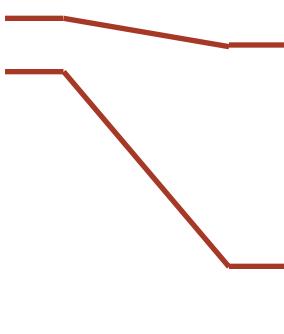
- Correlator manager
 - Deploy cluster
 - Process configuration files
 - Launch correlation
- Cluster manager
 - Methods for starting/stopping cluster
 - Methods for filesystem initialization
- Task launcher
 - Methods for correlation job launch
- Map
 - VDIF read and corner-turning
 - Initial integer-sample alignment
- Reduce
 - Fractional sample shift
 - Fringe rotation
 - Fourier transform
 - Fractional Sample Correction
 - Multiply and accumulate



CorrelX Implementation Details

Project	Language	Lines of code
CorrelX	Python	9k 2k comments
Libs.	Python	2M
Hadoop	Java	2M
Lustre	C	1M

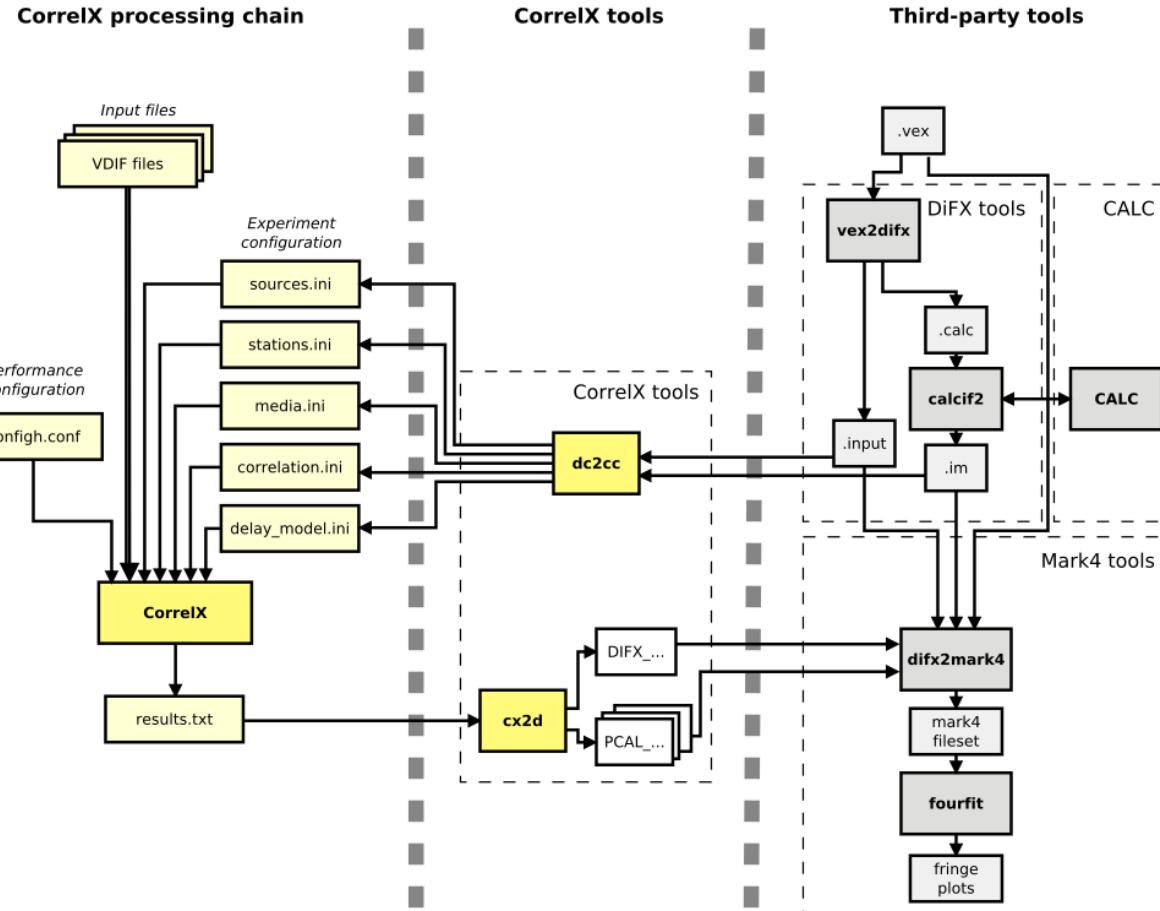
This is code maintained by
open-source projects



CorrelX Layer	LOC (approx.)
Launcher and FS Management	2k
Application (map, reduce)	1.8k
Libraries (VDIF, quant., FX...)	4k
Tools (dc2cc, cx2d...)	1.2k

This is the code the VLBI
community needs to maintain

Tool Chain Integration

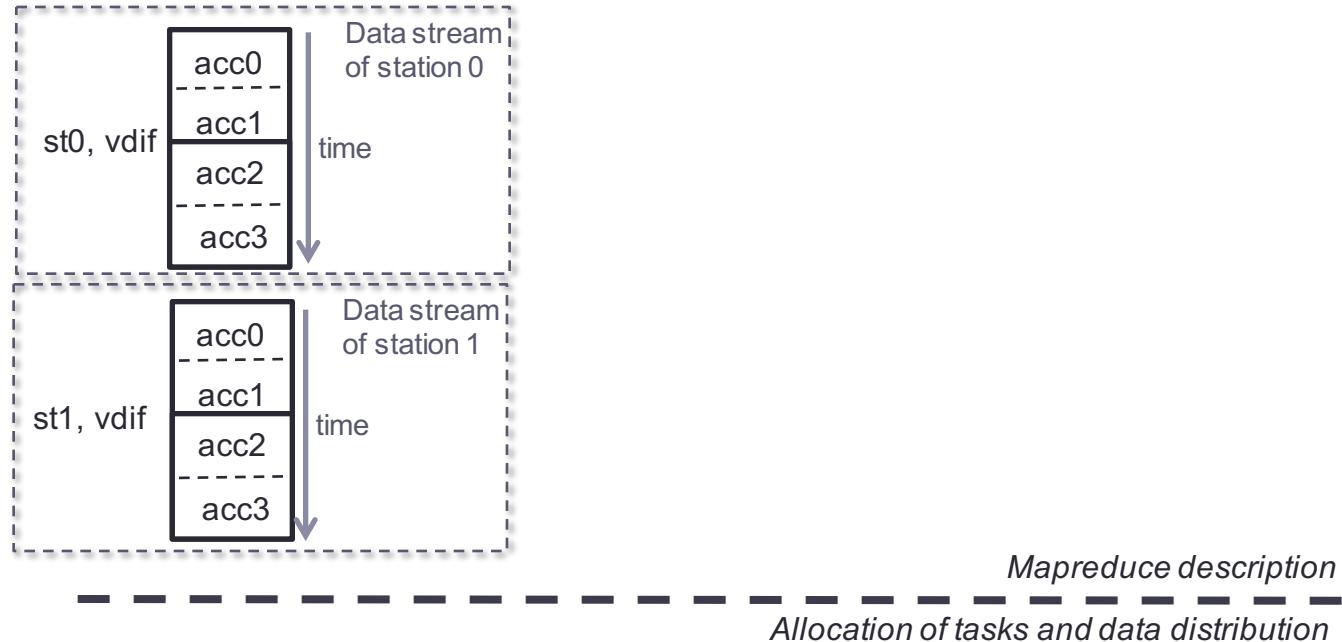


- Tools for integration with existing processing chains:
 - Input file converter from .input and .im
 - Output file converter to SWIN and PCAL files
- Integration with Mark4 post-processing chain
 - dc2cc: configuration
 - cx2d: output conversion
- Experiment files
 - sources.ini
 - stations.ini
 - media.ini
 - correlation.ini
 - delay_model.ini
- Correlator configuration
 - configh.conf

Correlation Example

- Correlation example:

- Cluster:
 - 3 nodes.
 - 2 cores/node.
- Experiment:
 - 2 stations.
 - 2 files, 2 parts/file
 $\Rightarrow 2 \times 2 = 4$ mappers
 - 2 bands.
 - 4 accumulation periods
 $\Rightarrow 2 \times 4 = 8$ reducers

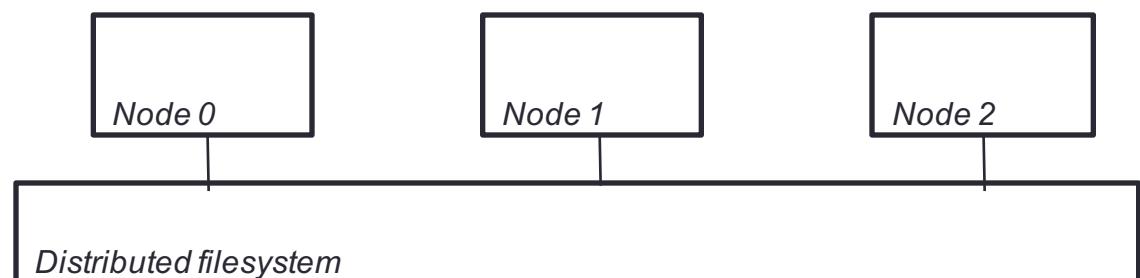


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

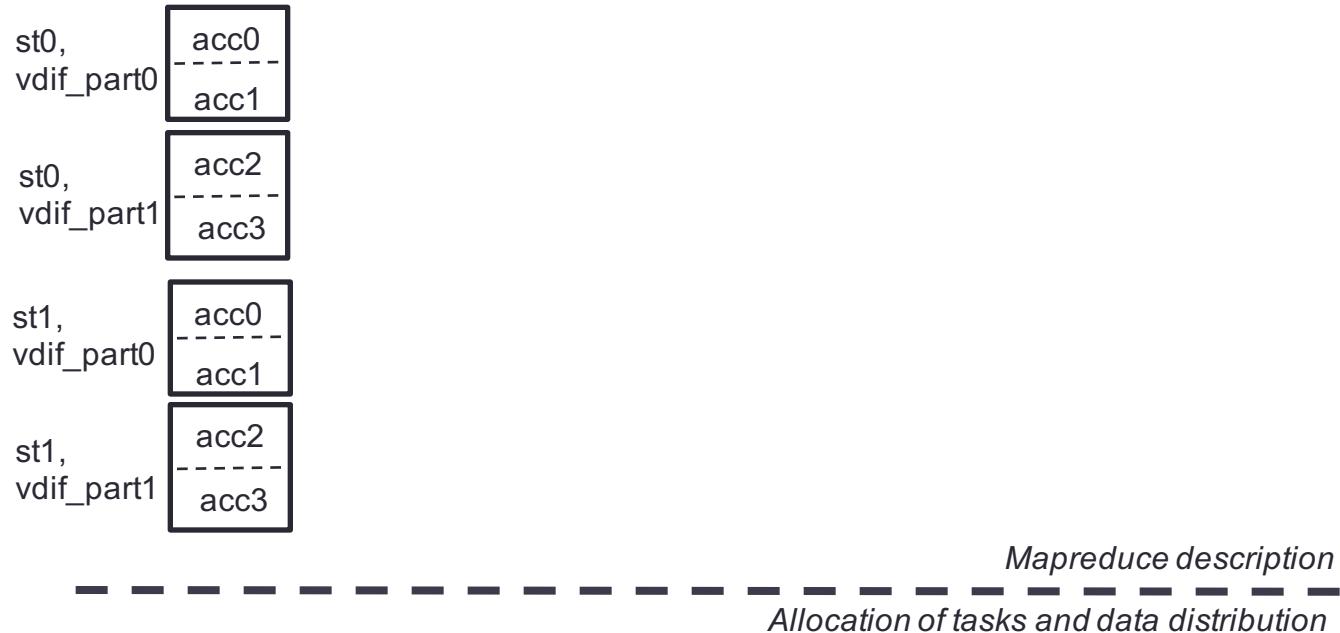
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.

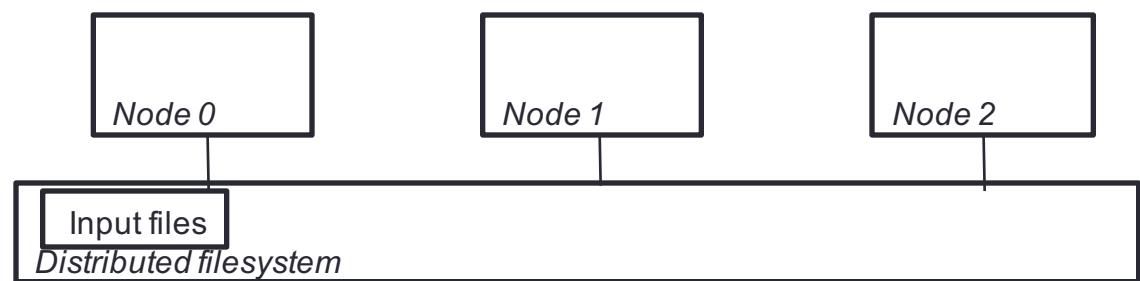


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

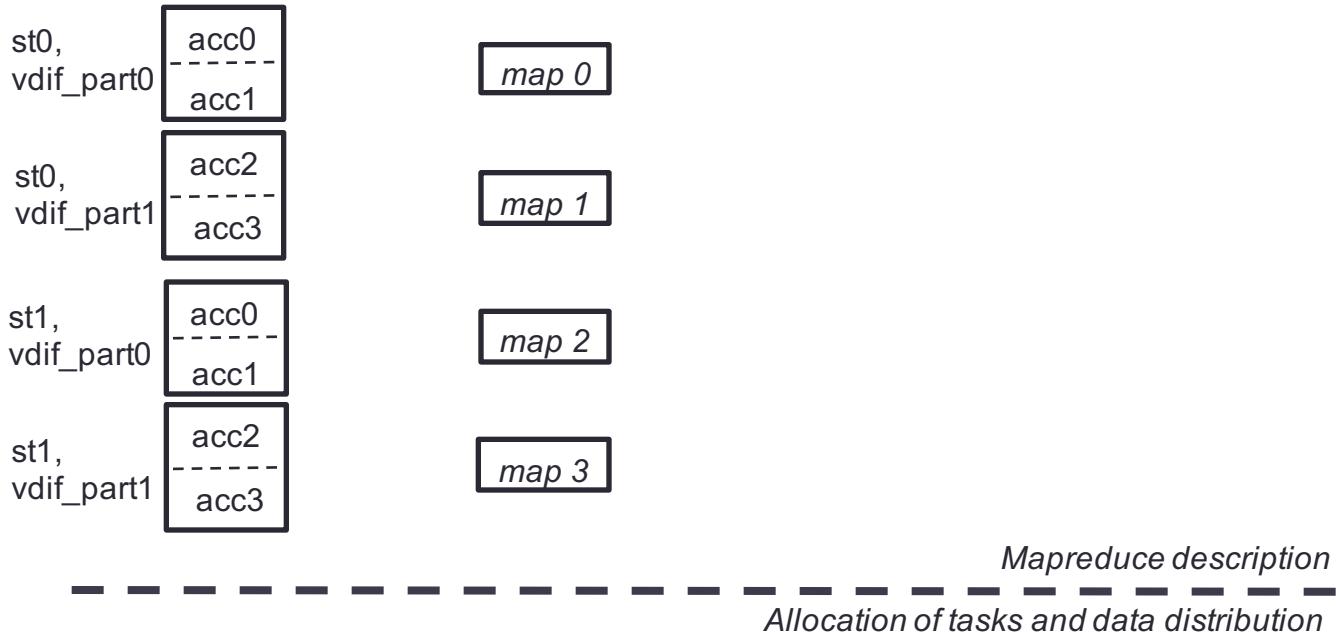
Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.

2. The framework launches one map per inputfile.

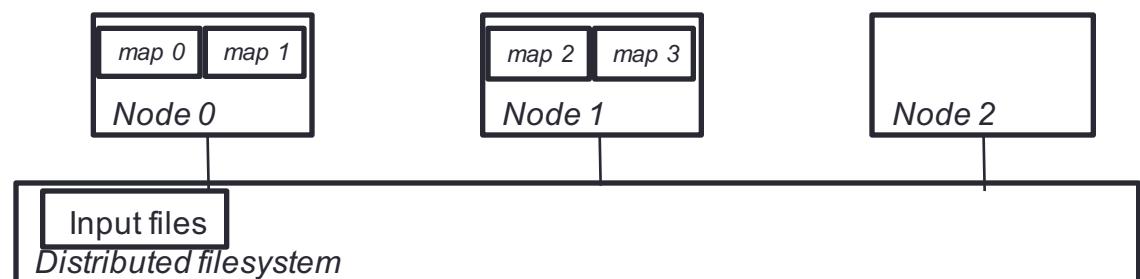


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

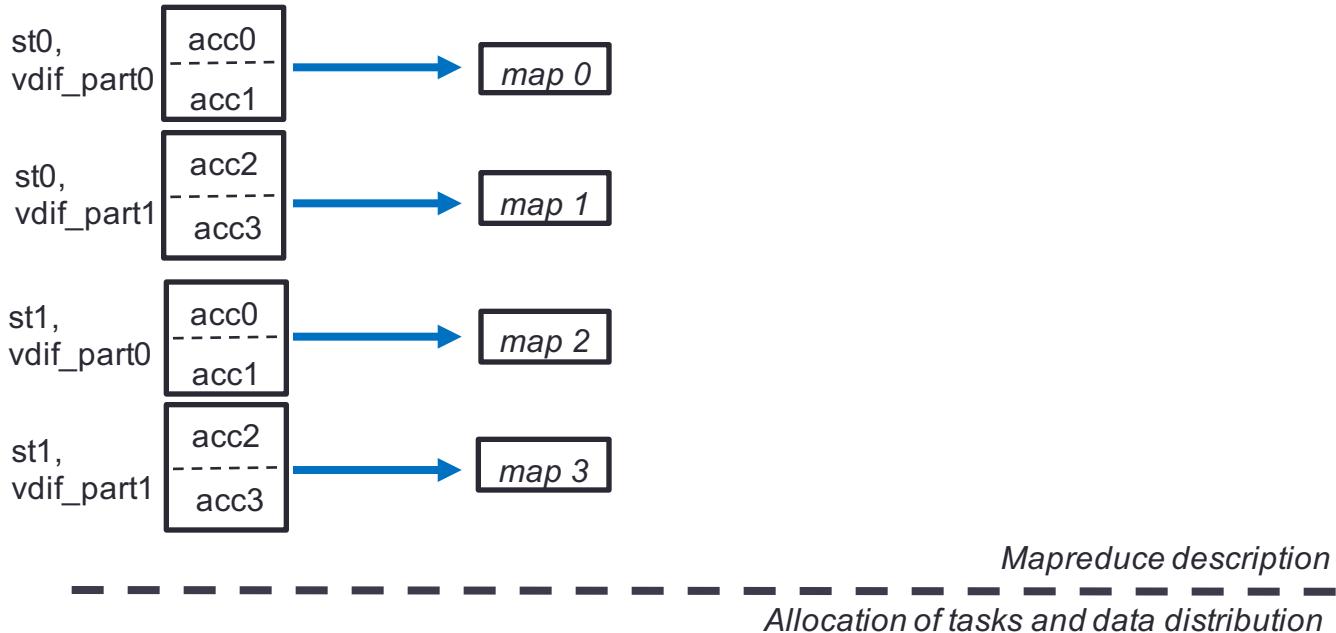
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.

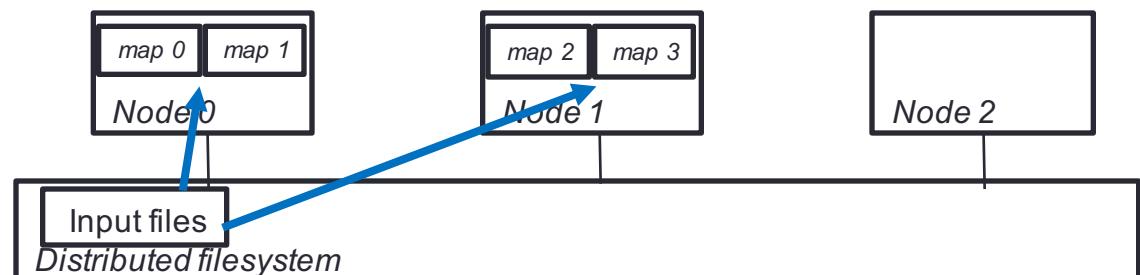


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

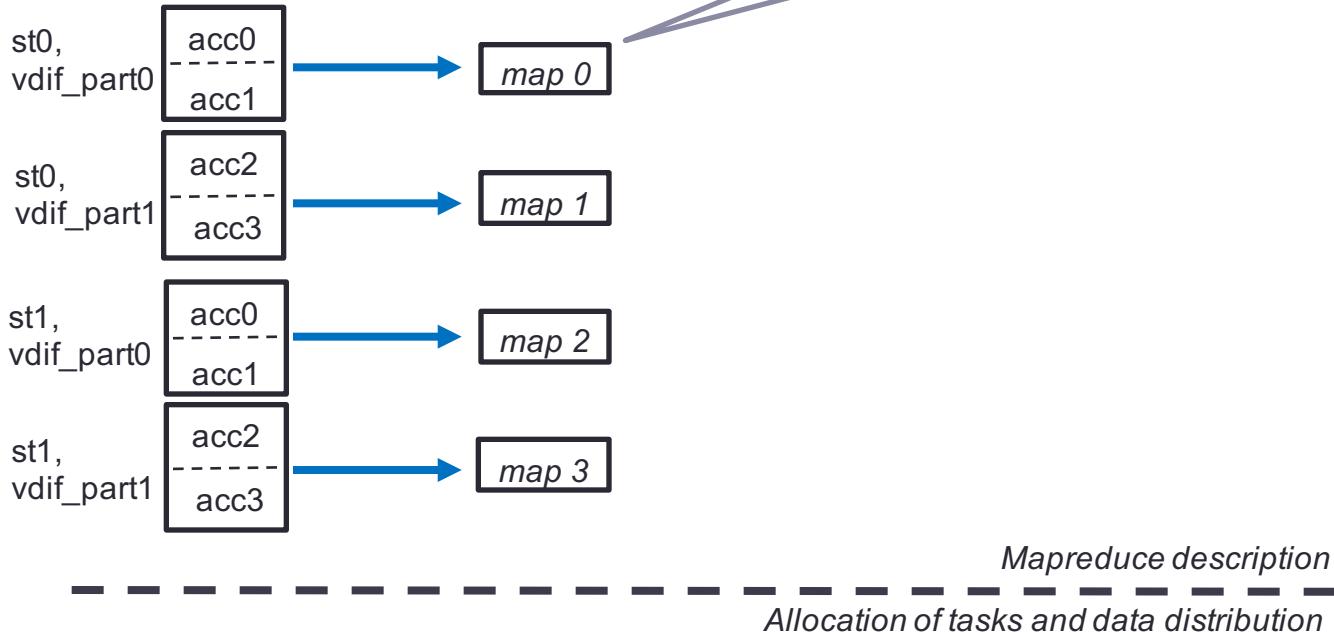
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.

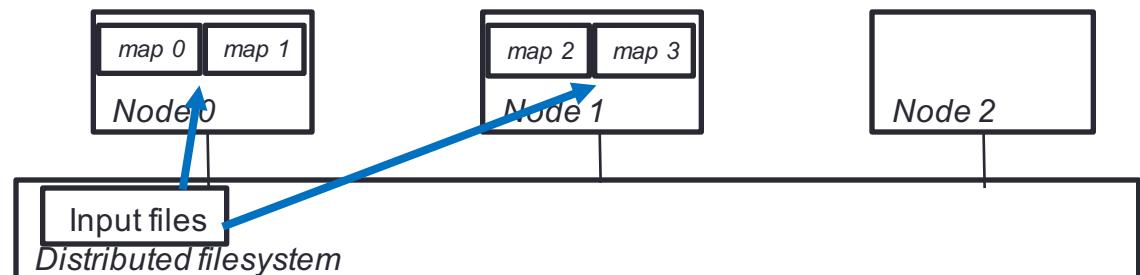


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

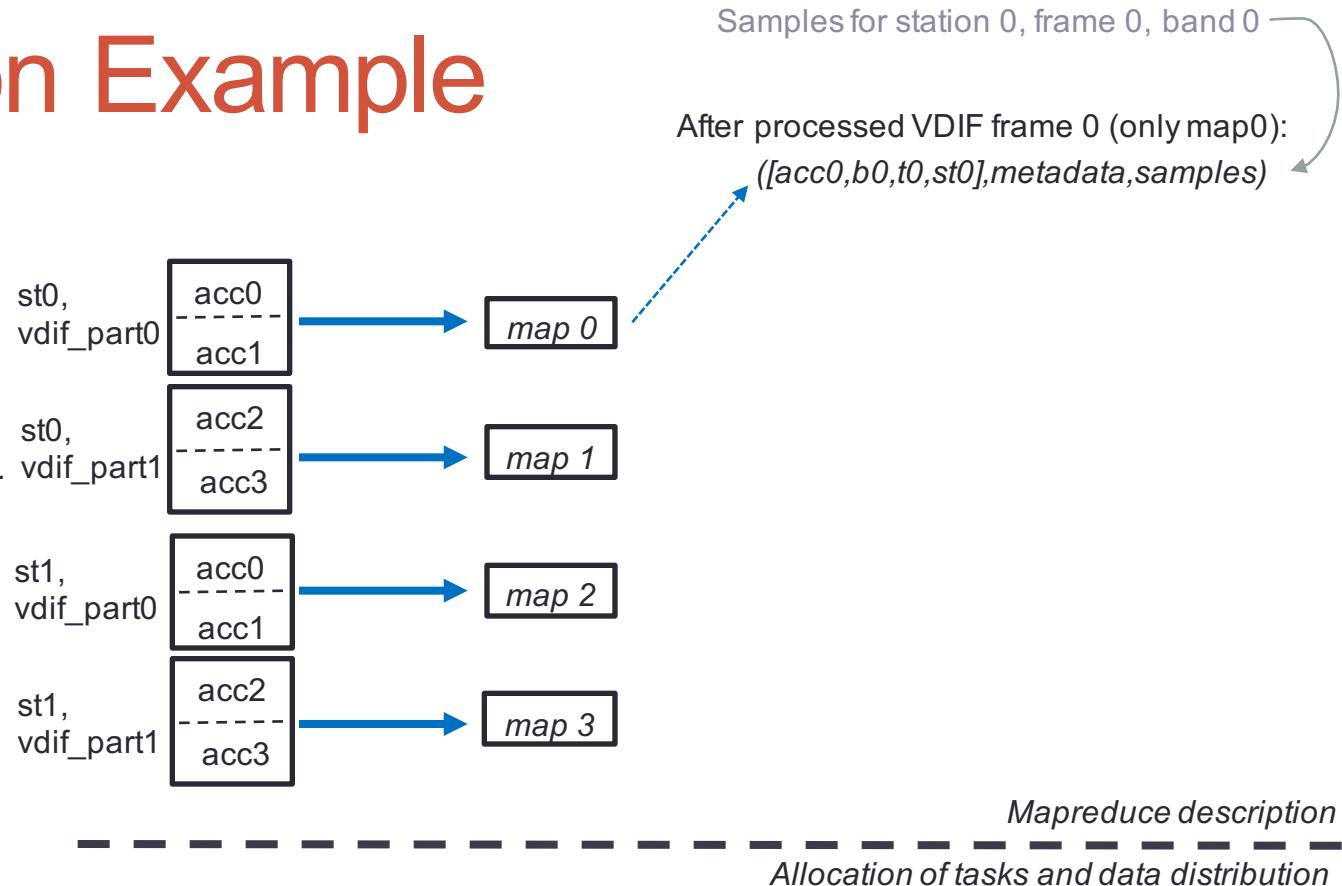
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

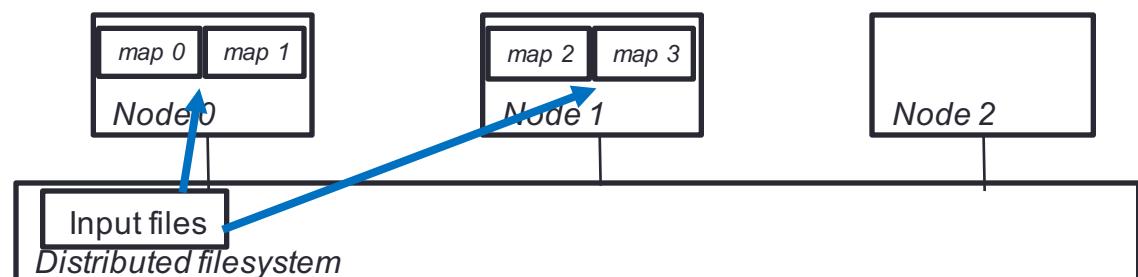


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

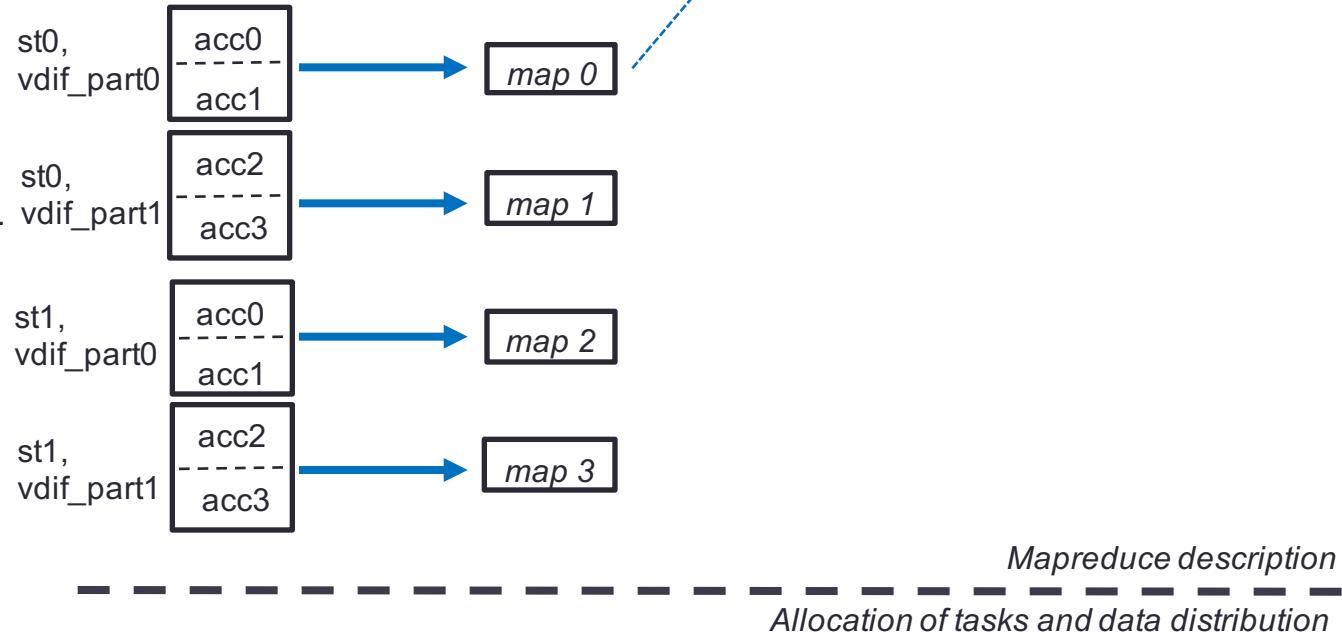
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

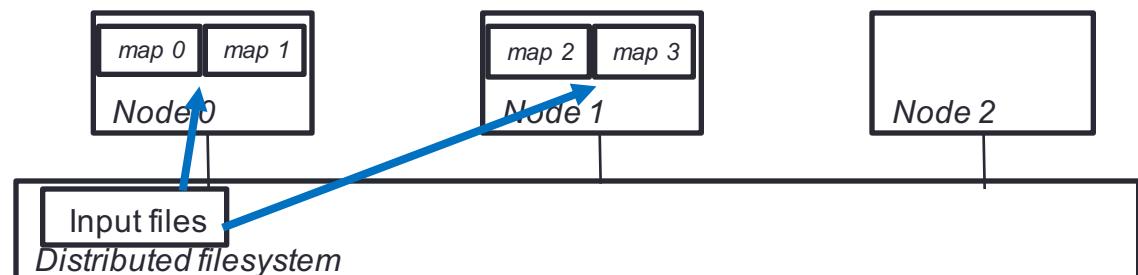


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

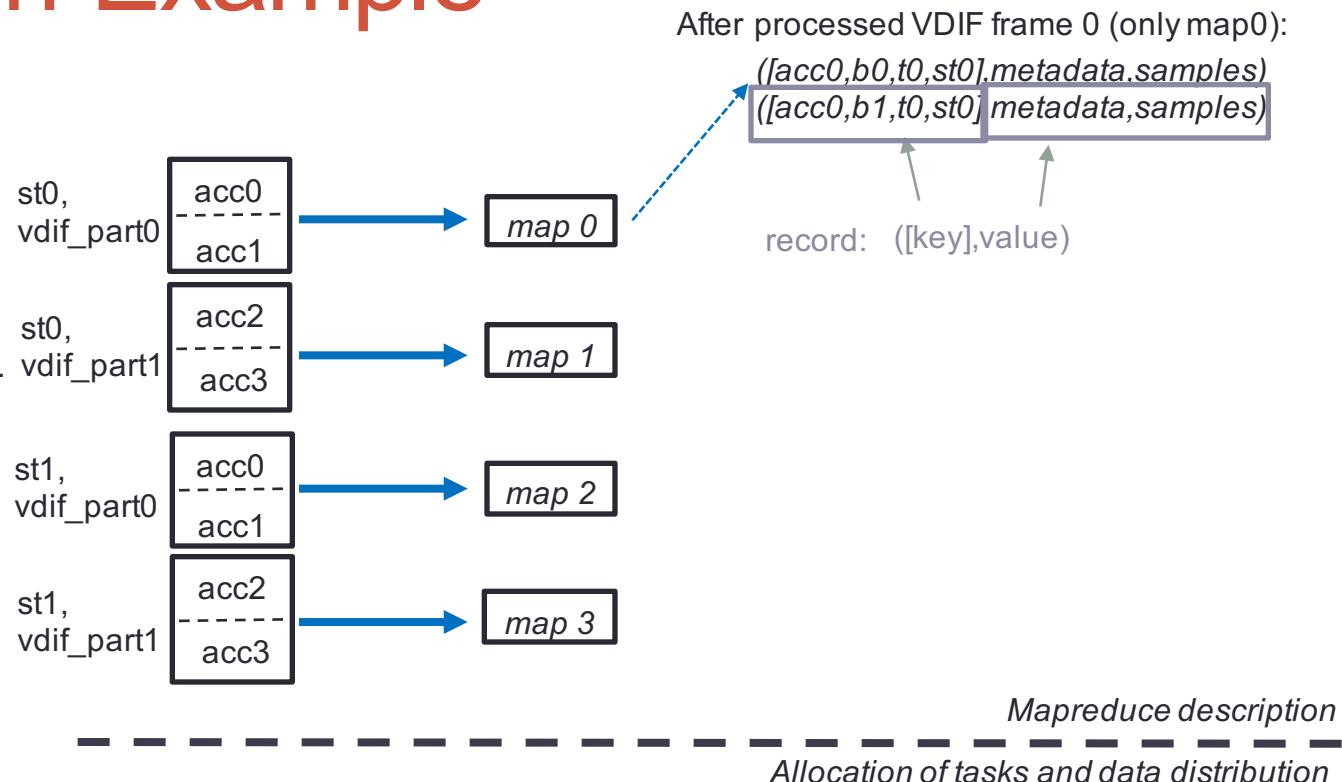
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

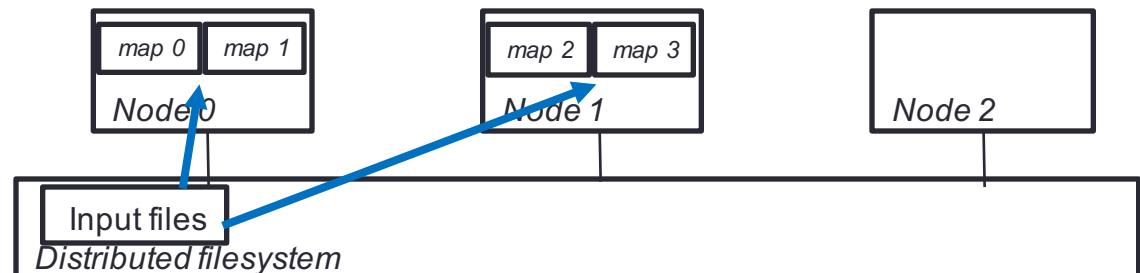


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

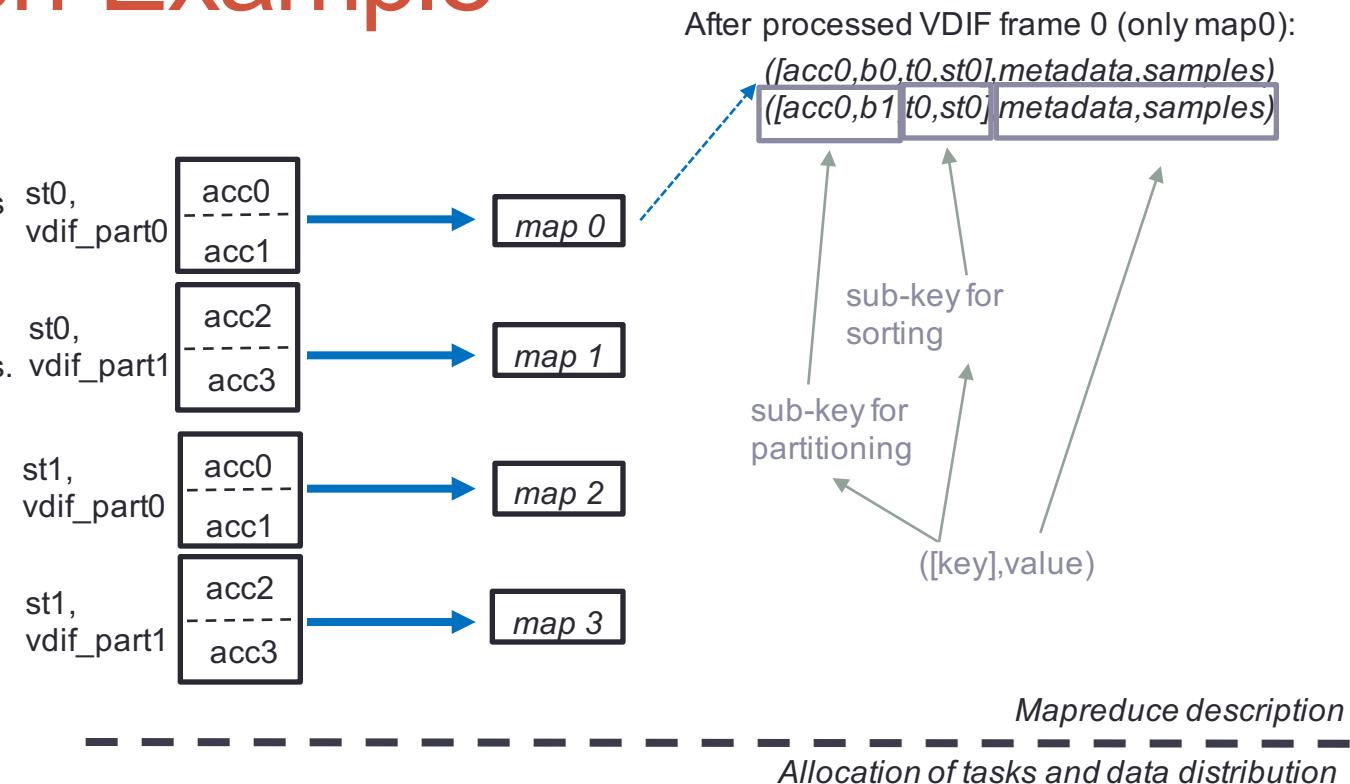
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: [sts]*[splits]=4
Number of reducers: [accs]*[bands]=8



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

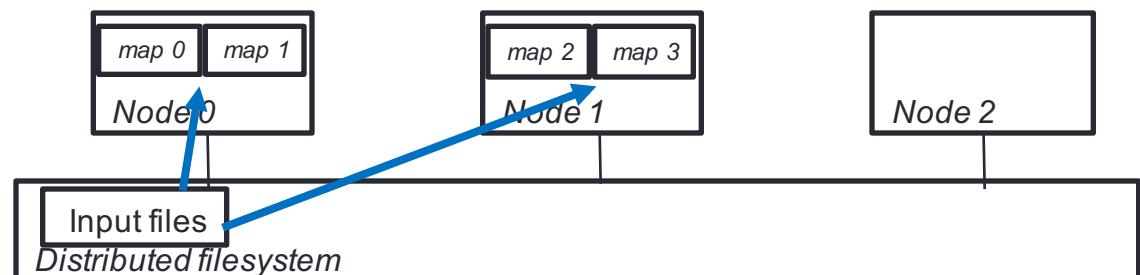


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

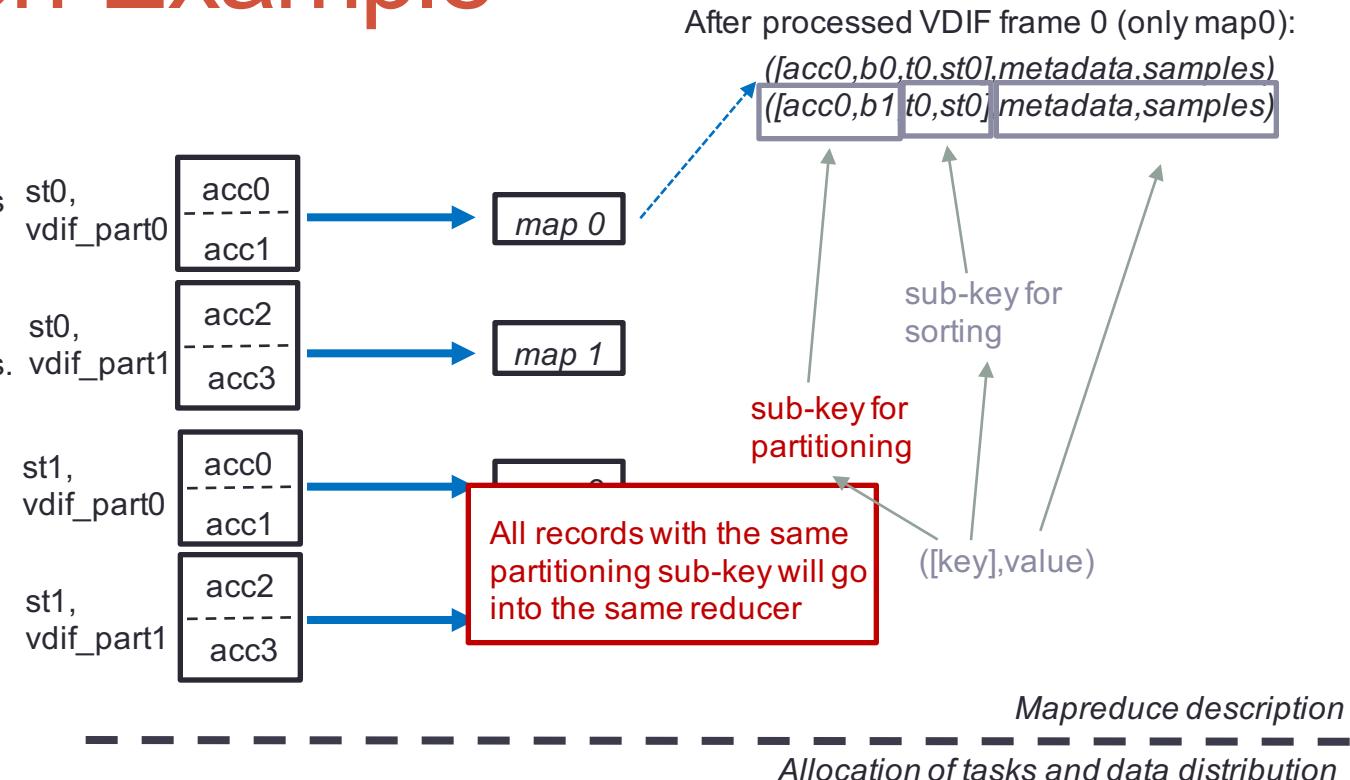
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

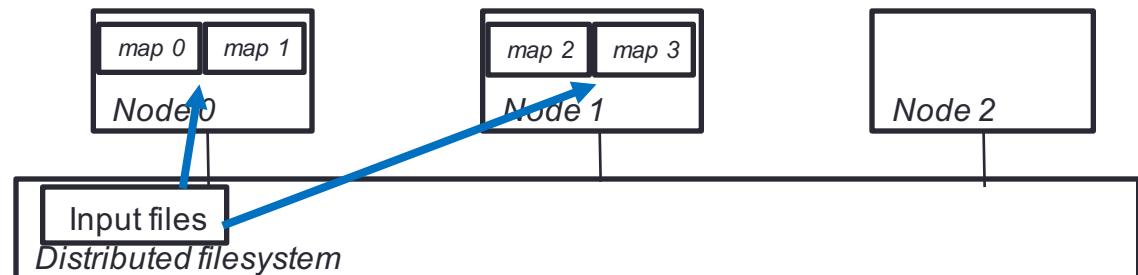


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

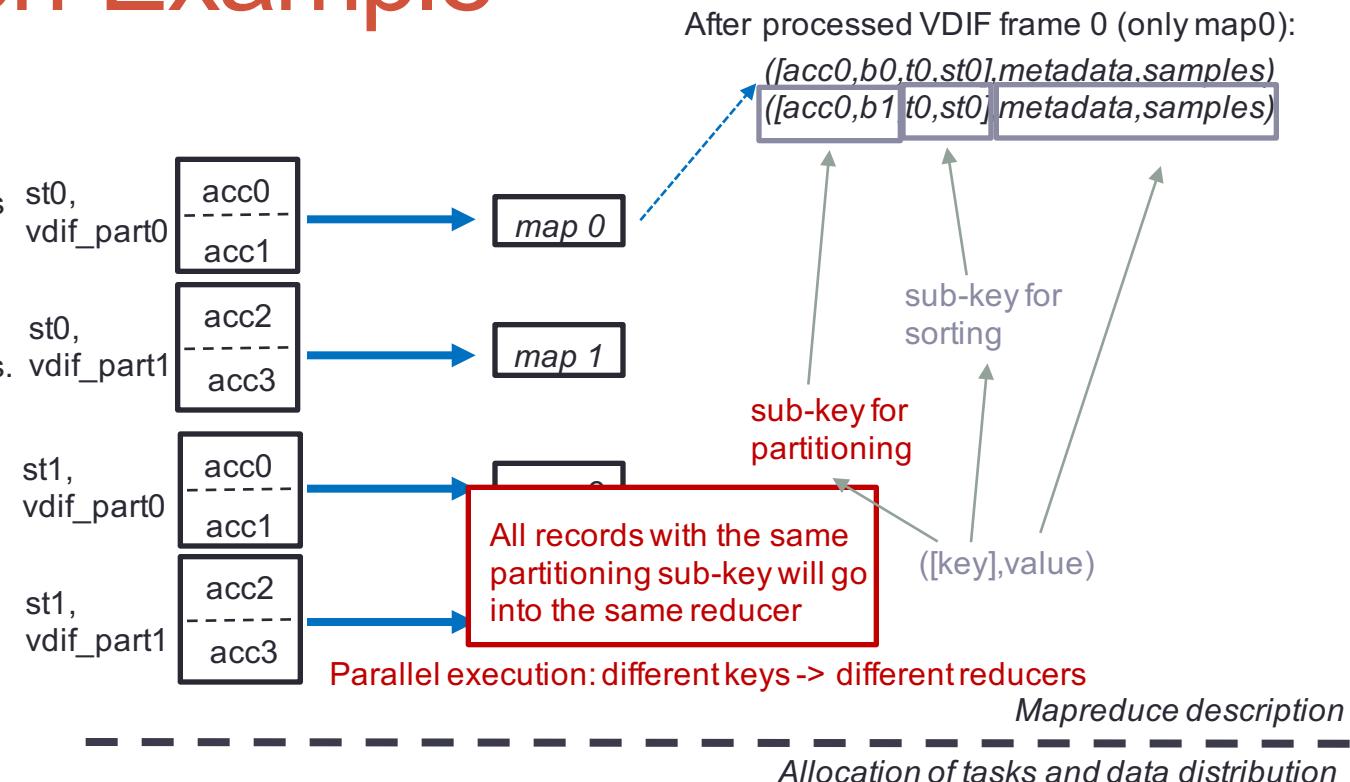
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

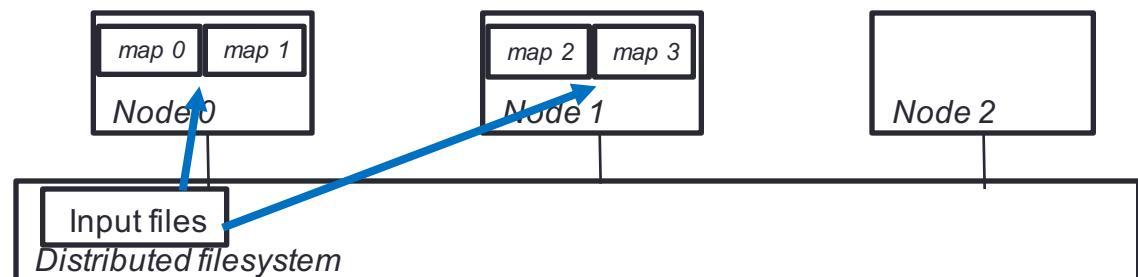


MapReduce correlation example:

Number of nodes: 3
 Number of cores per node: 2

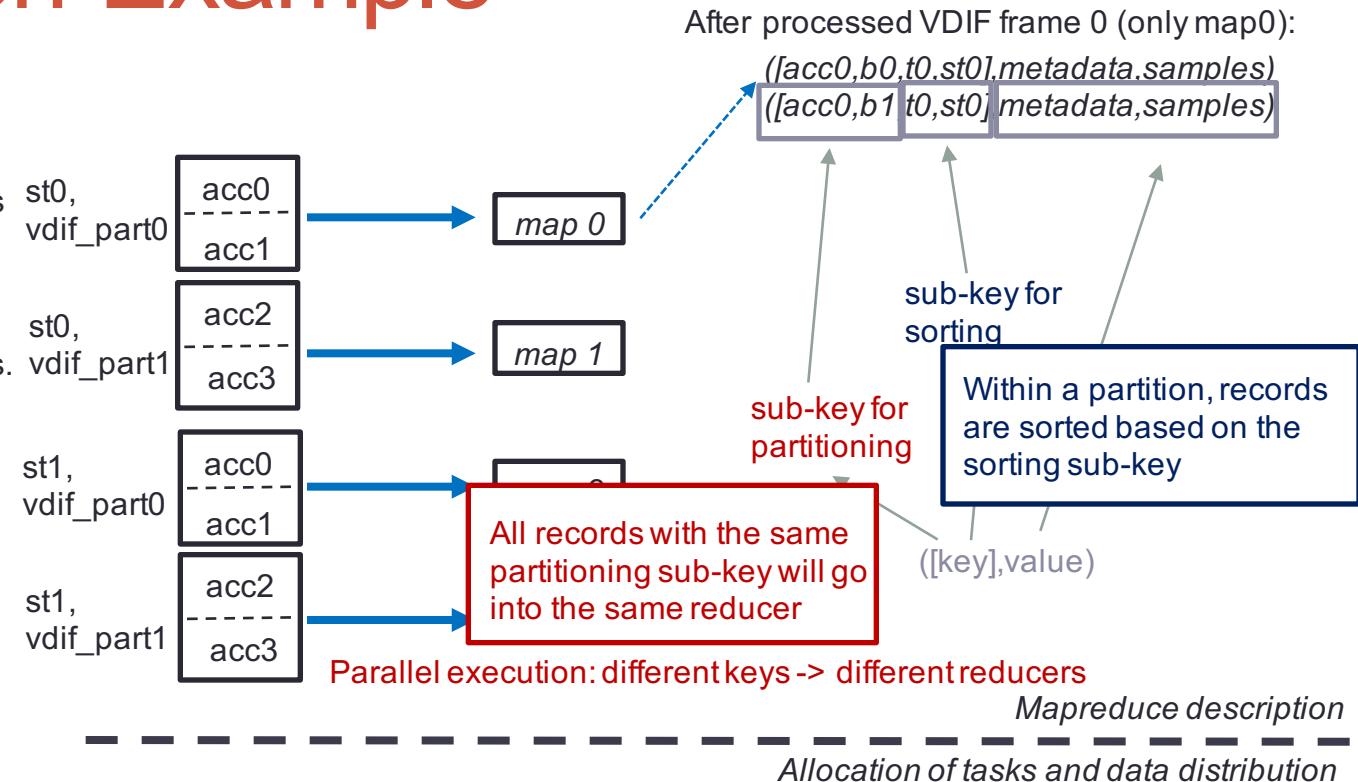
Stations: 2
 Input splits: 2 per file.
 Acc. Periods: 4.
 Bands: 2.

Number of mappers: $[sts]*[splits]=4$
 Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

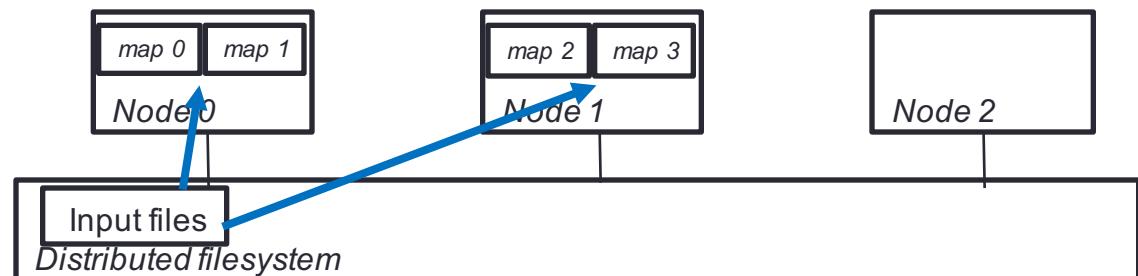


MapReduce correlation example:

Number of nodes: 3
 Number of cores per node: 2

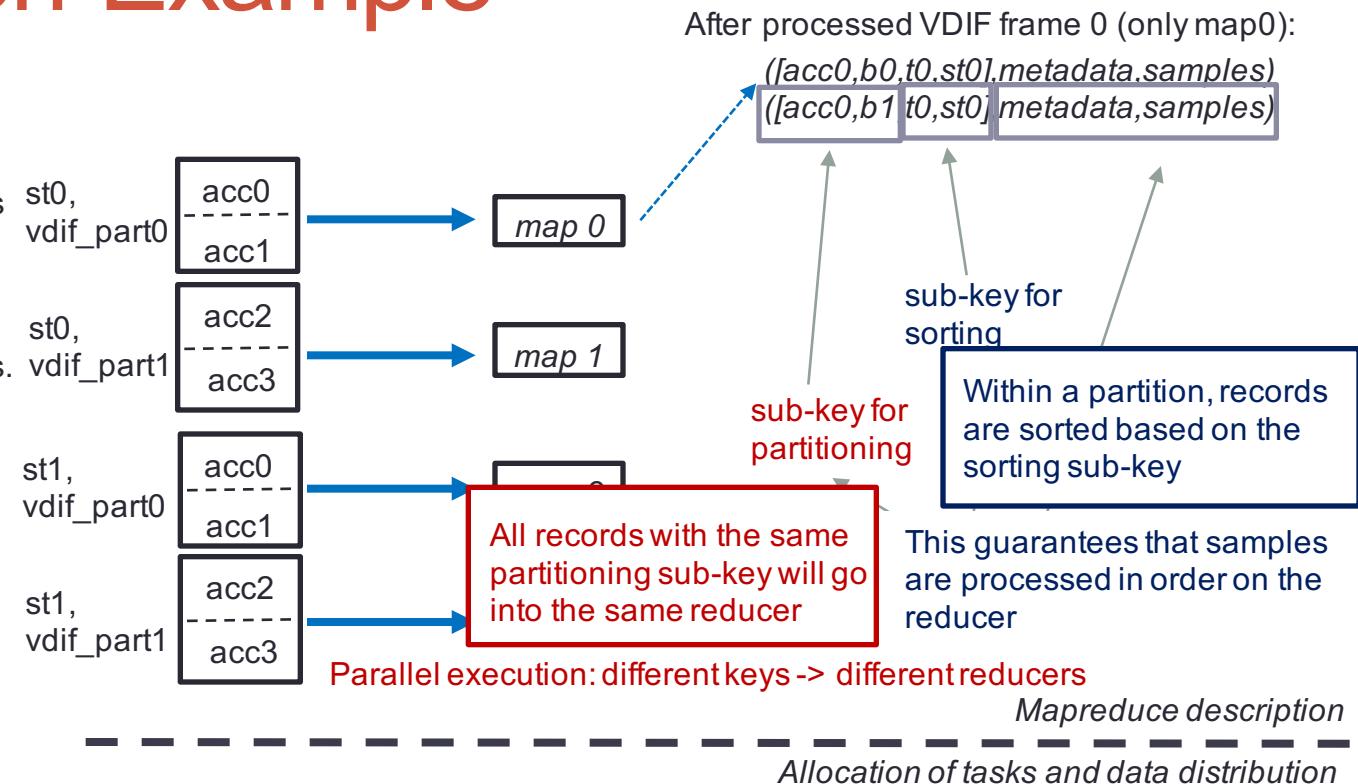
Stations: 2
 Input splits: 2 per file.
 Acc. Periods: 4.
 Bands: 2.

Number of mappers: $[sts]*[splits]=4$
 Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

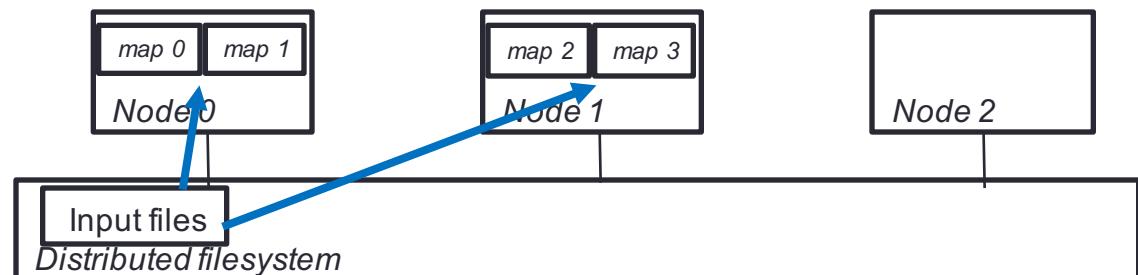


MapReduce correlation example:

Number of nodes: 3
 Number of cores per node: 2

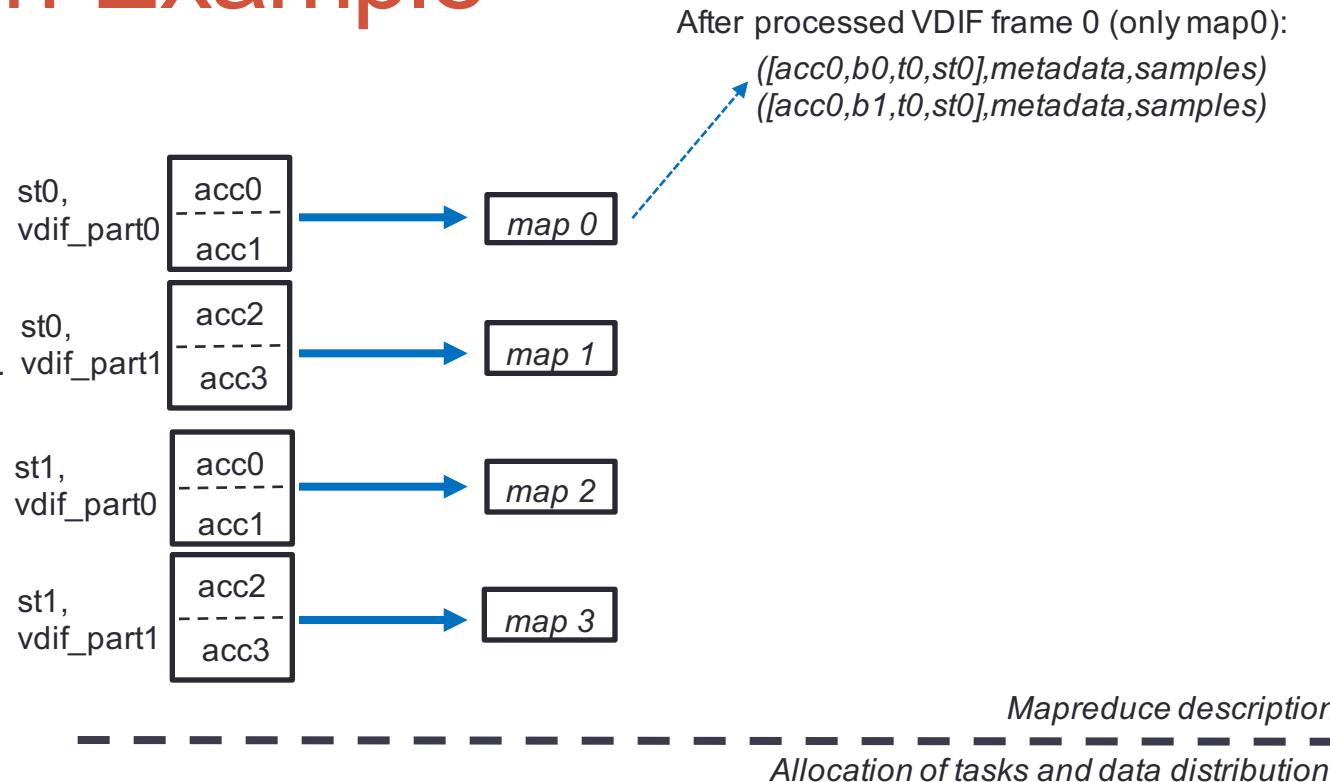
Stations: 2
 Input splits: 2 per file.
 Acc. Periods: 4.
 Bands: 2.

Number of mappers: $[sts]*[splits]=4$
 Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

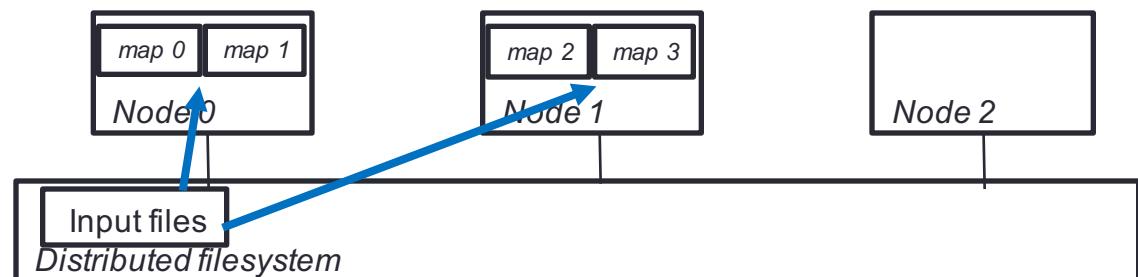


MapReduce correlation example:

Number of nodes: 3
 Number of cores per node: 2

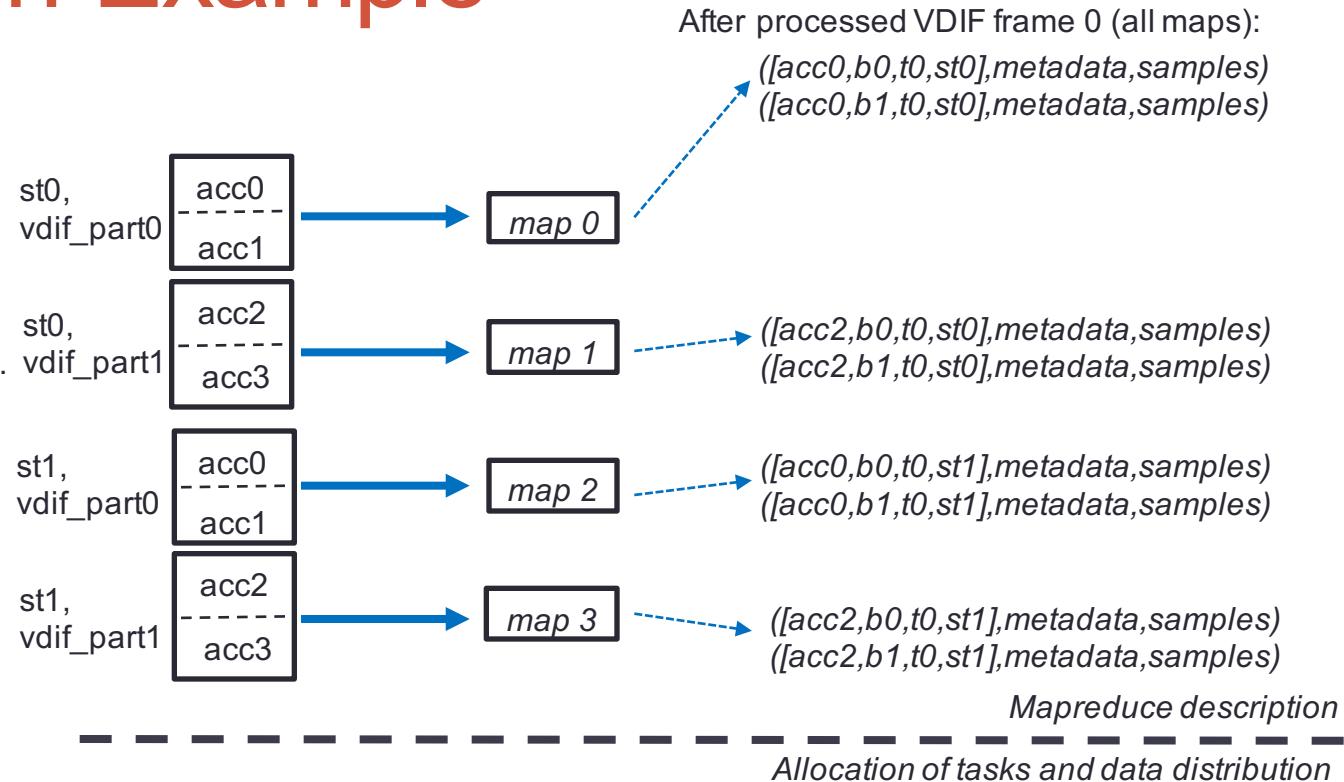
Stations: 2
 Input splits: 2 per file.
 Acc. Periods: 4.
 Bands: 2.

Number of mappers: $[sts]*[splits]=4$
 Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

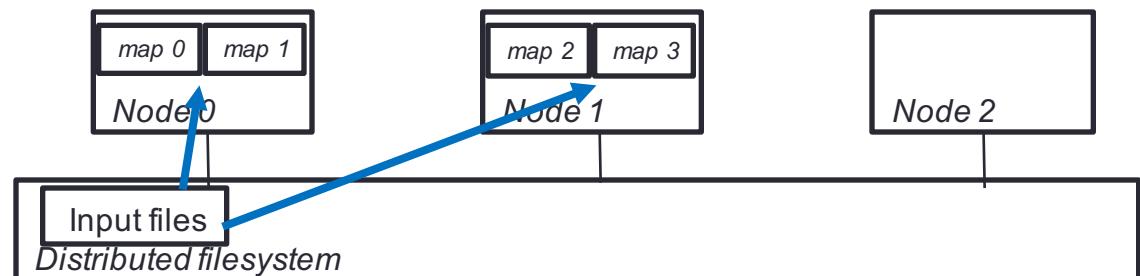


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

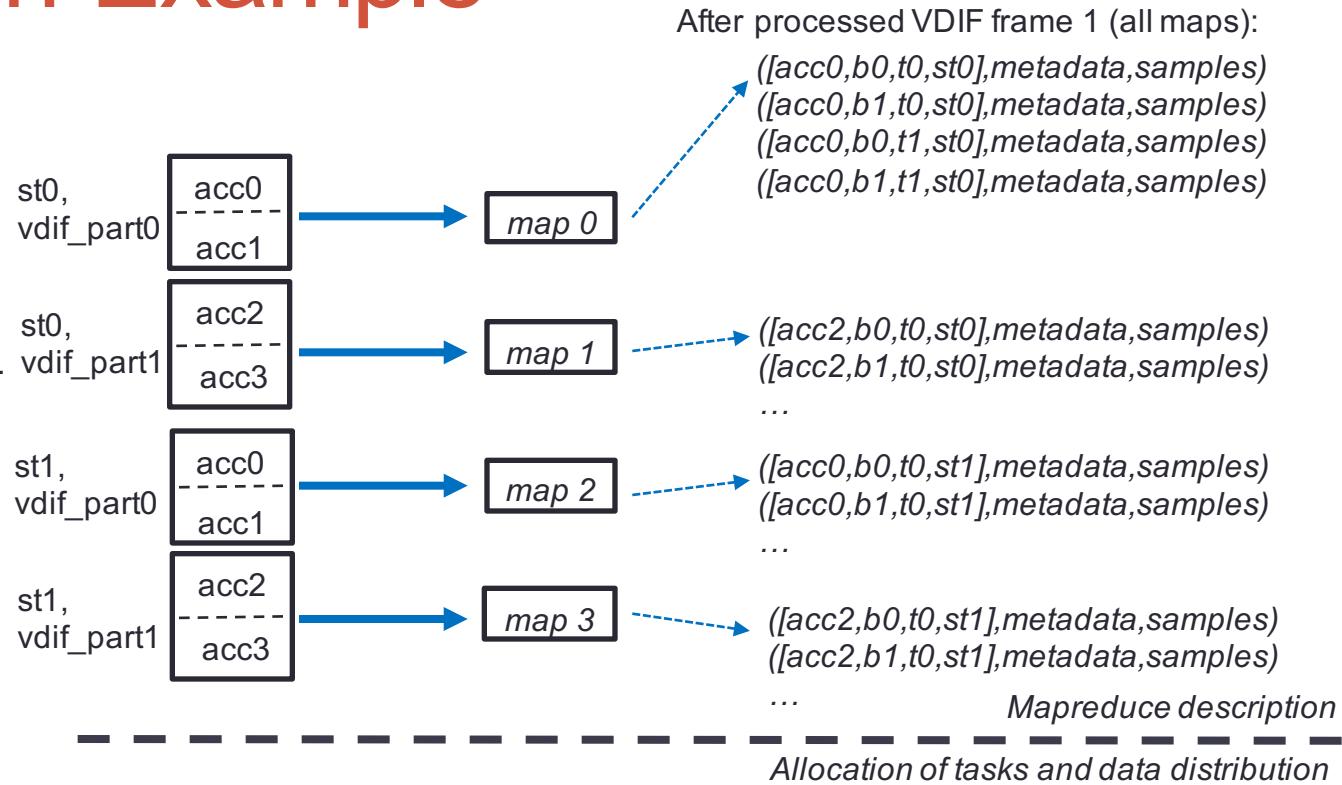
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

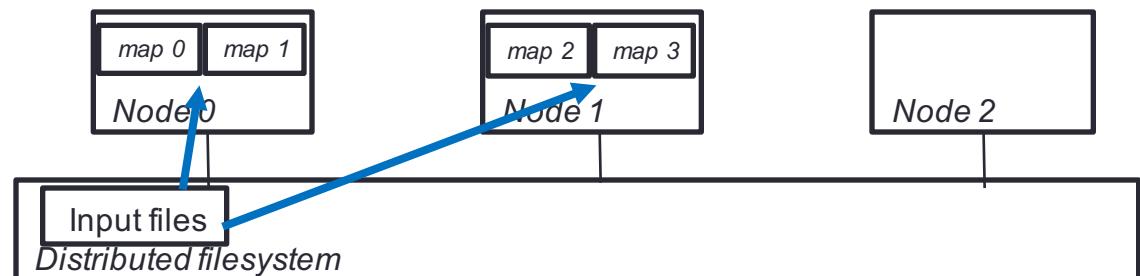


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

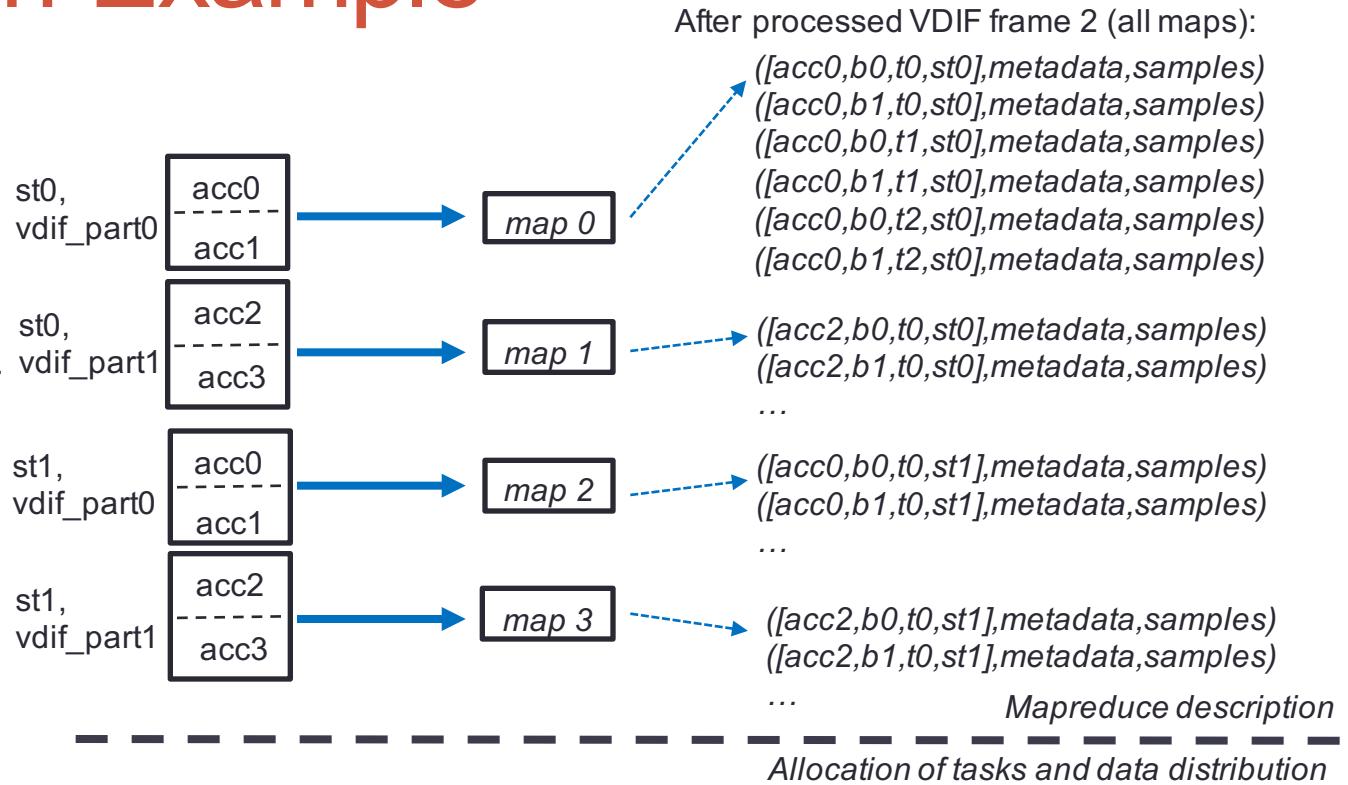
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

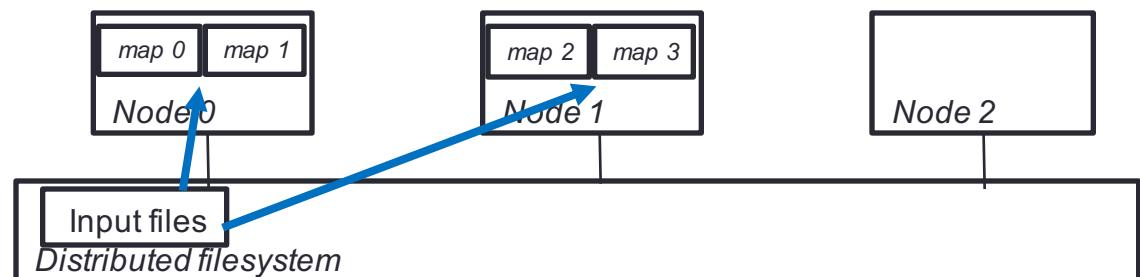


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

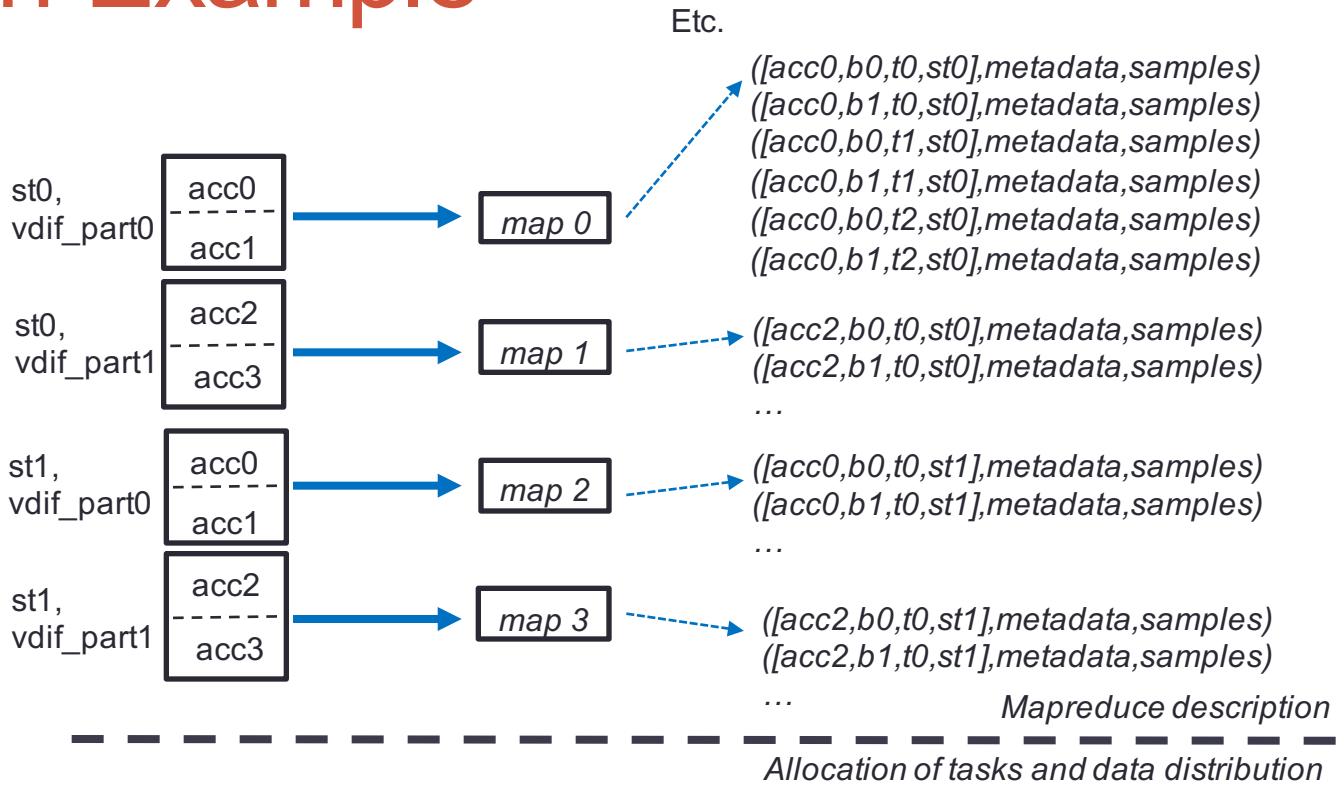
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.

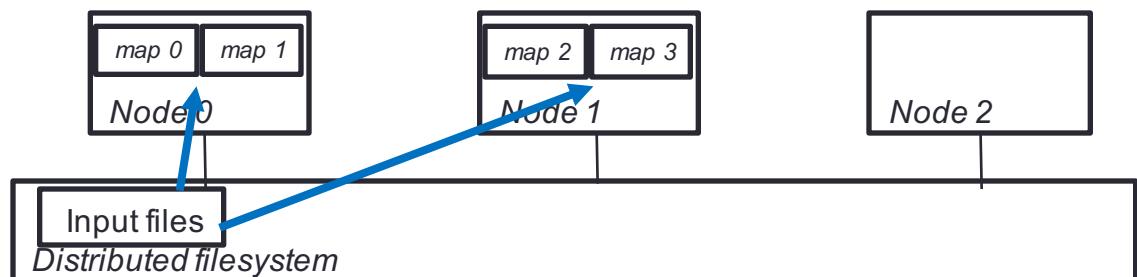


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

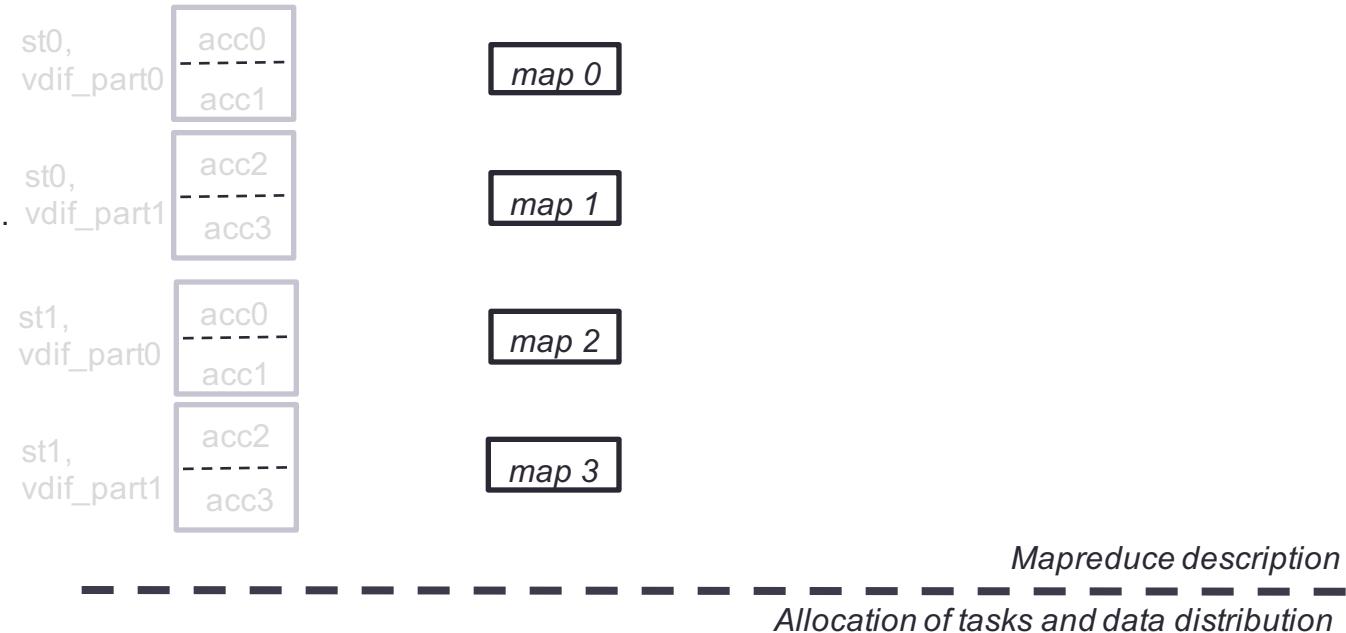
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.

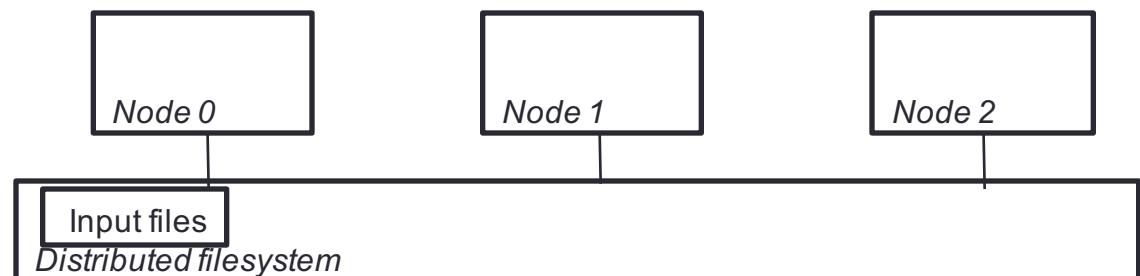


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

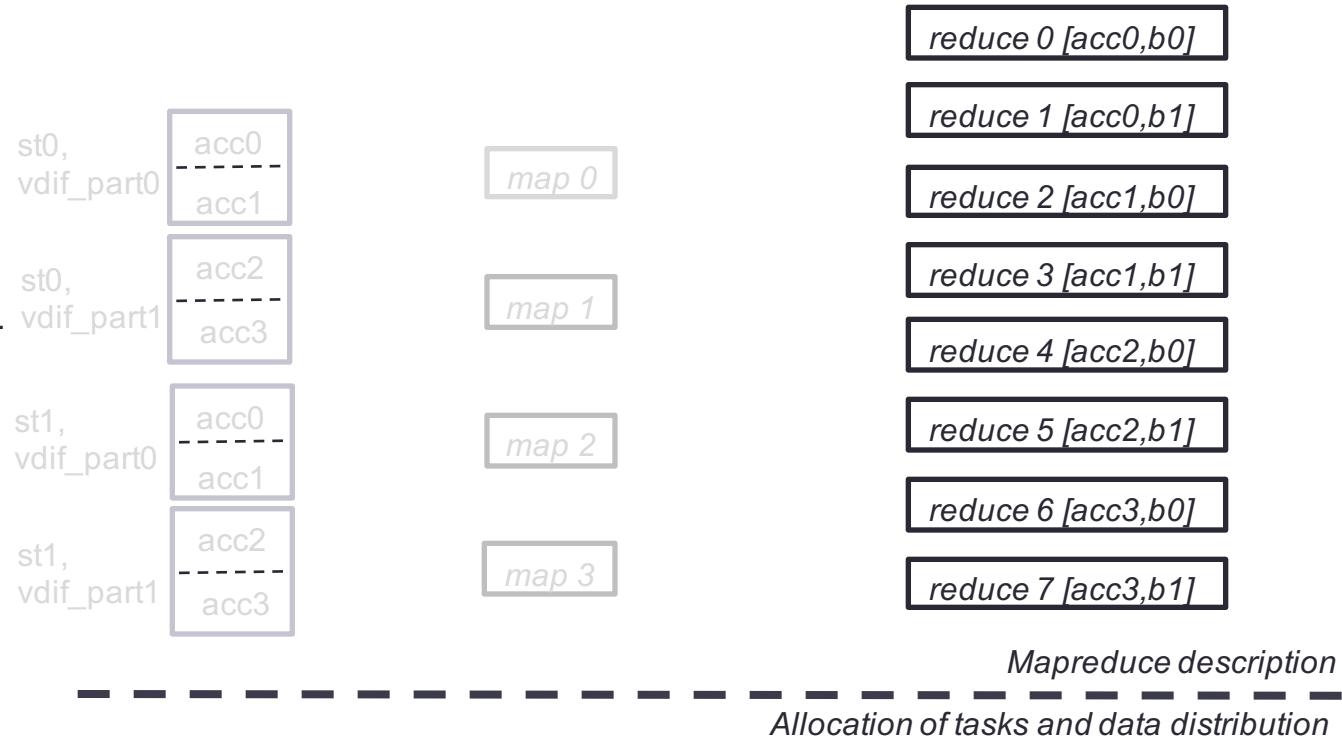
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.

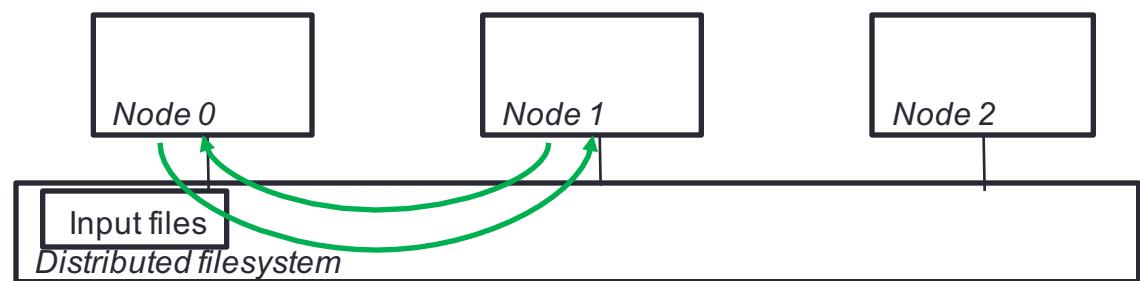


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

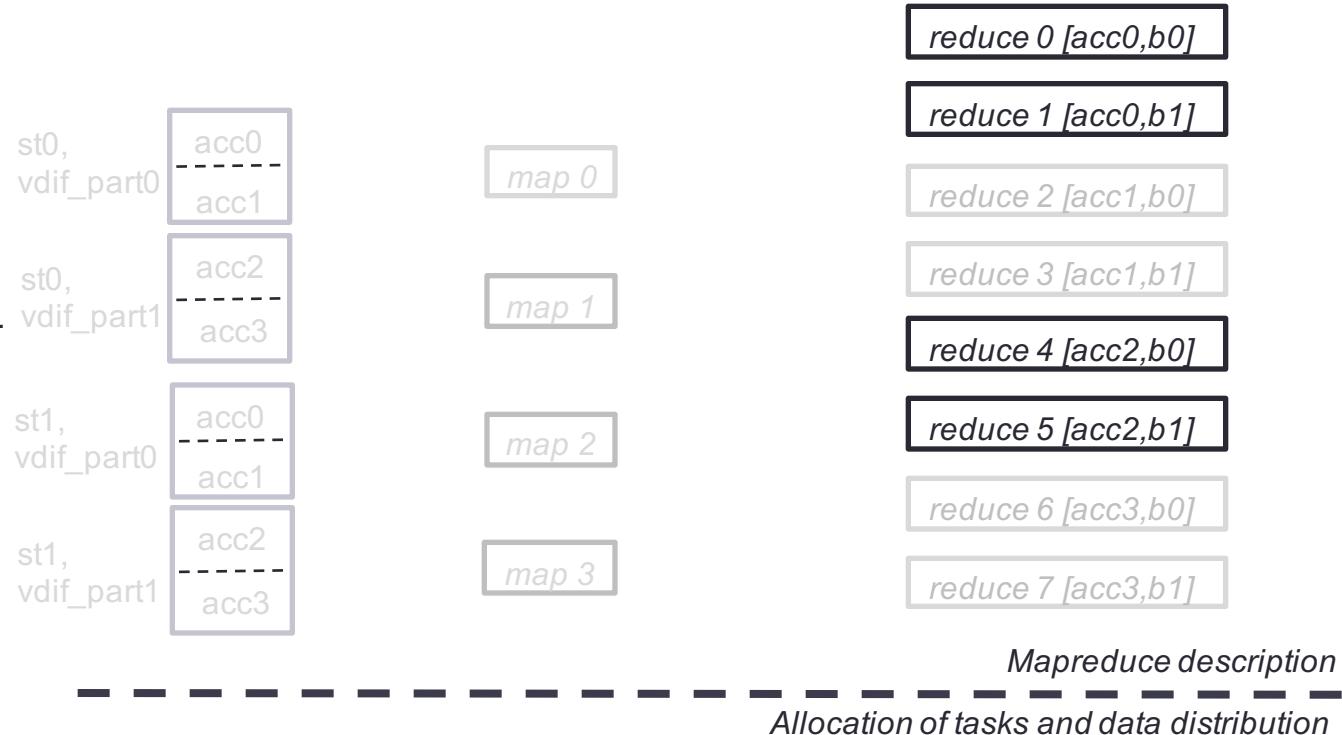
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.

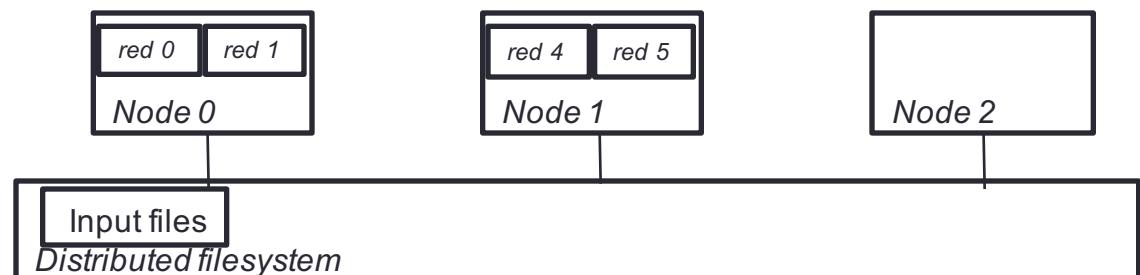


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

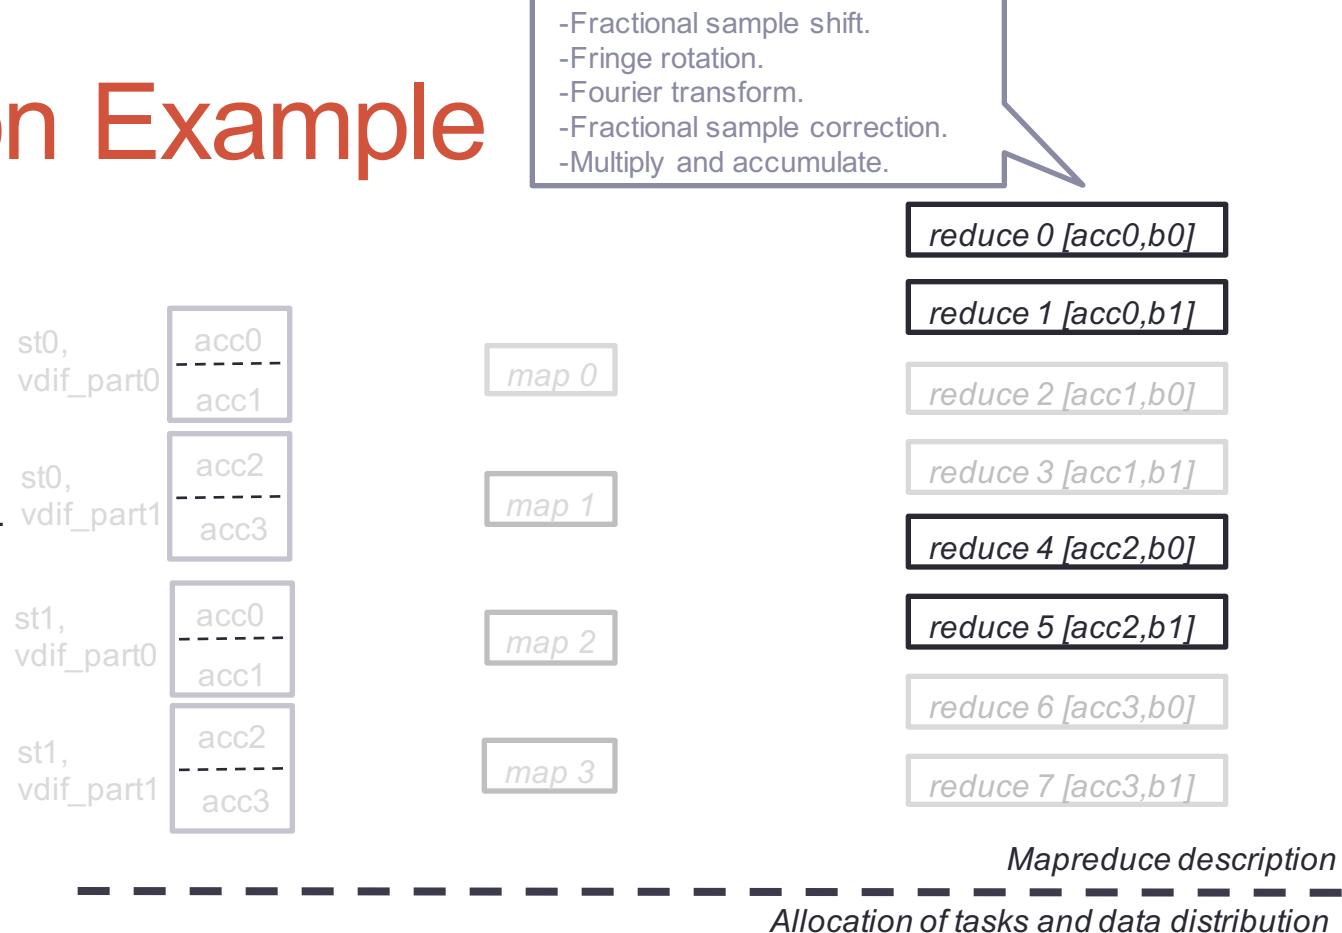
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per input file.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.

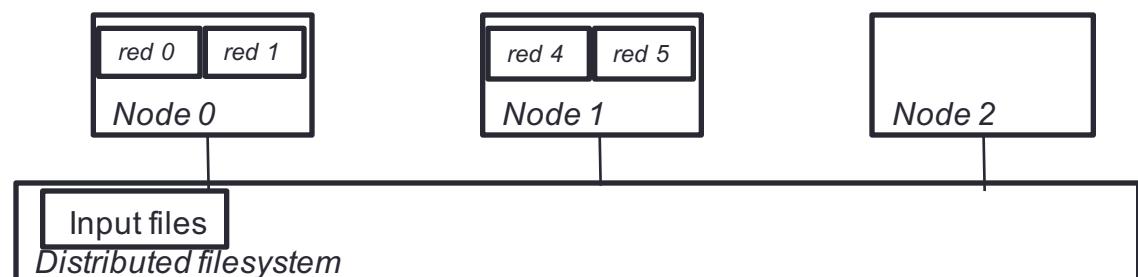


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

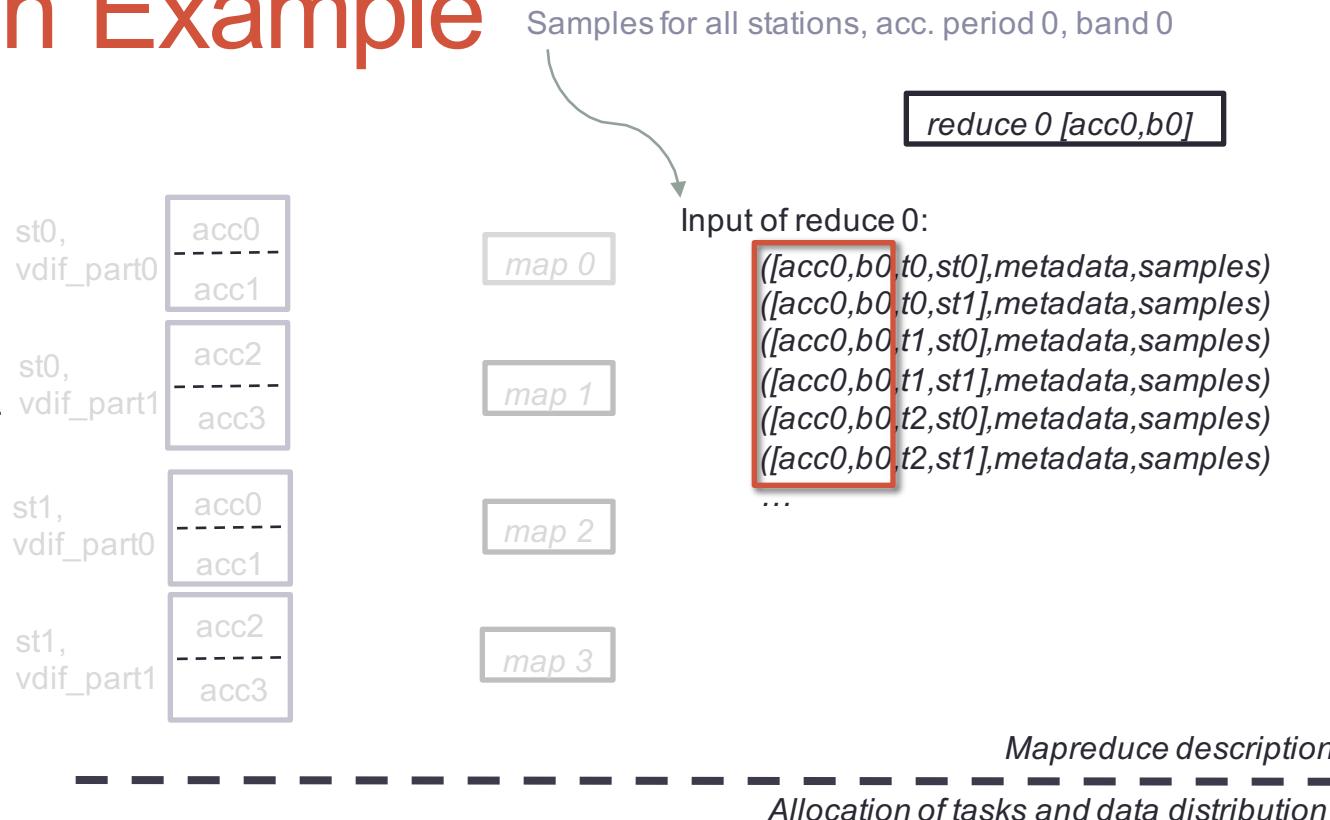
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.

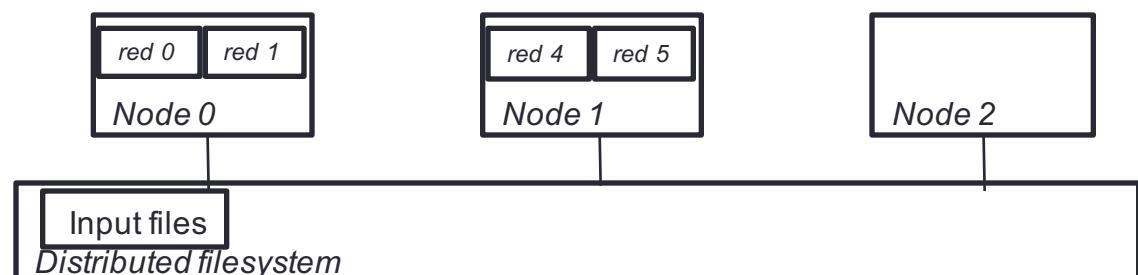


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

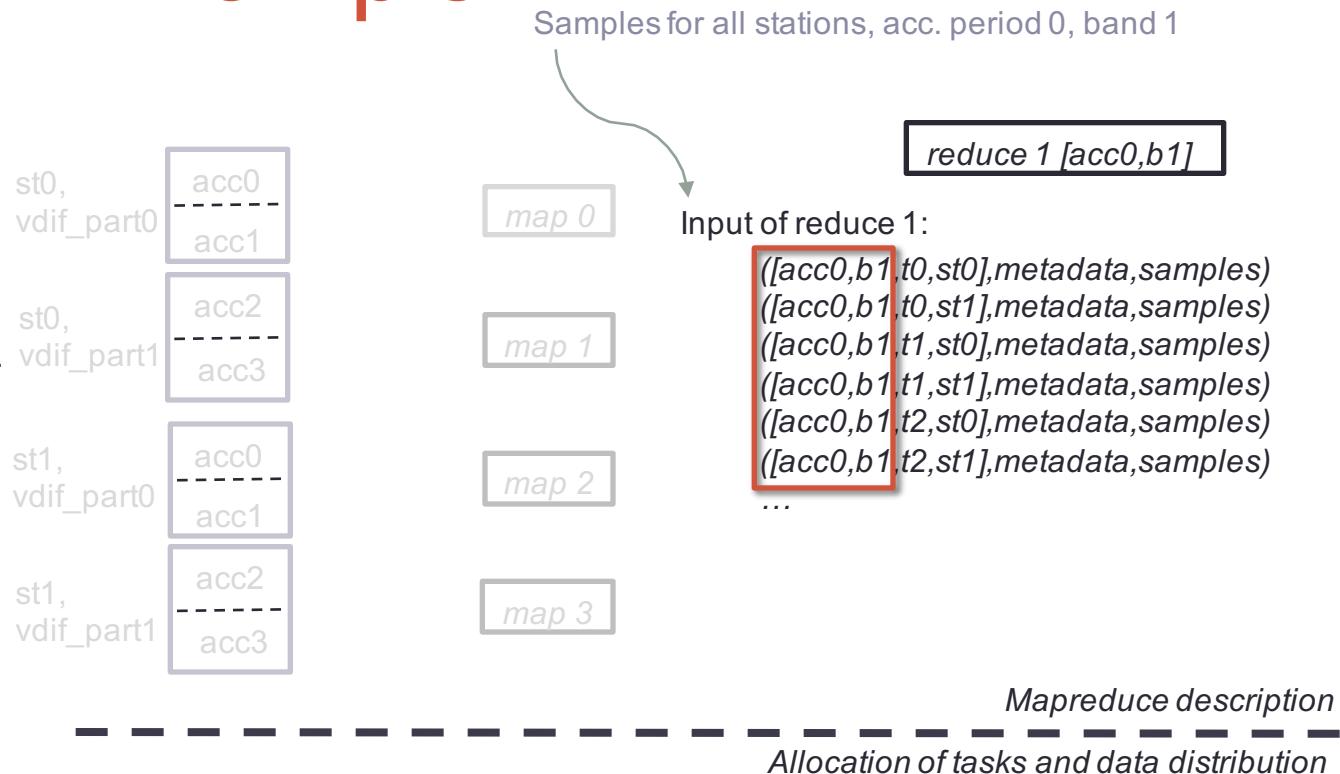
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.

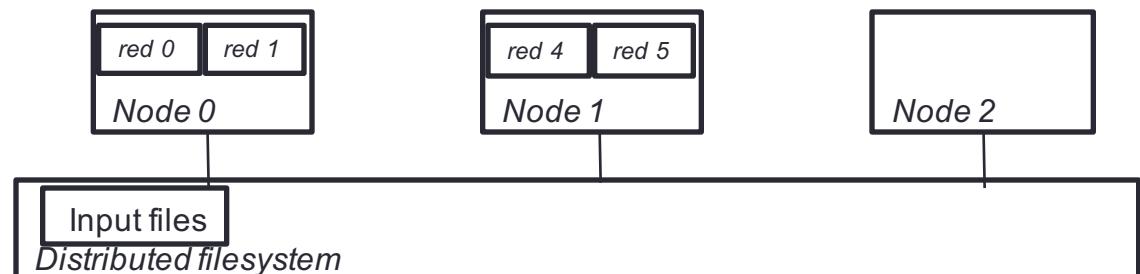


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

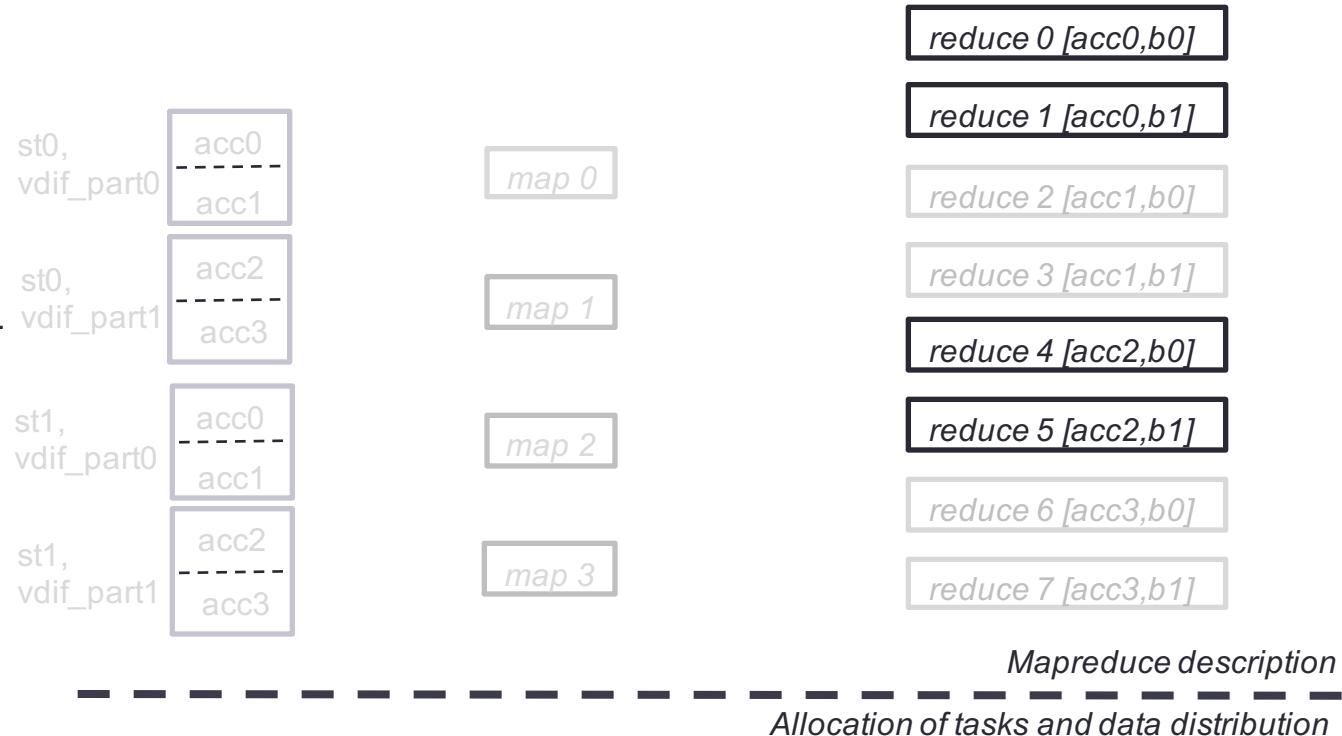
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.

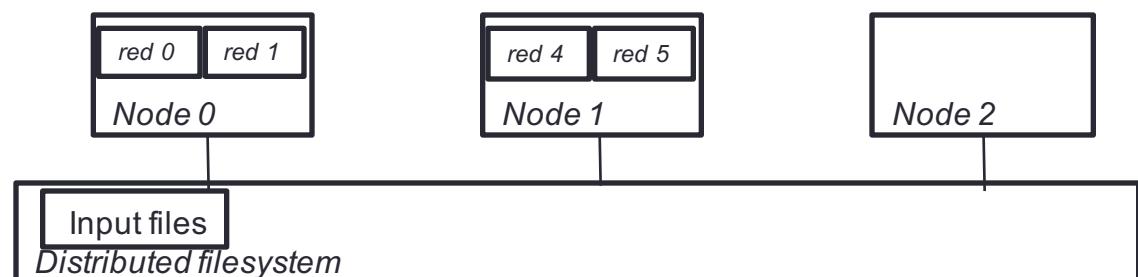


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

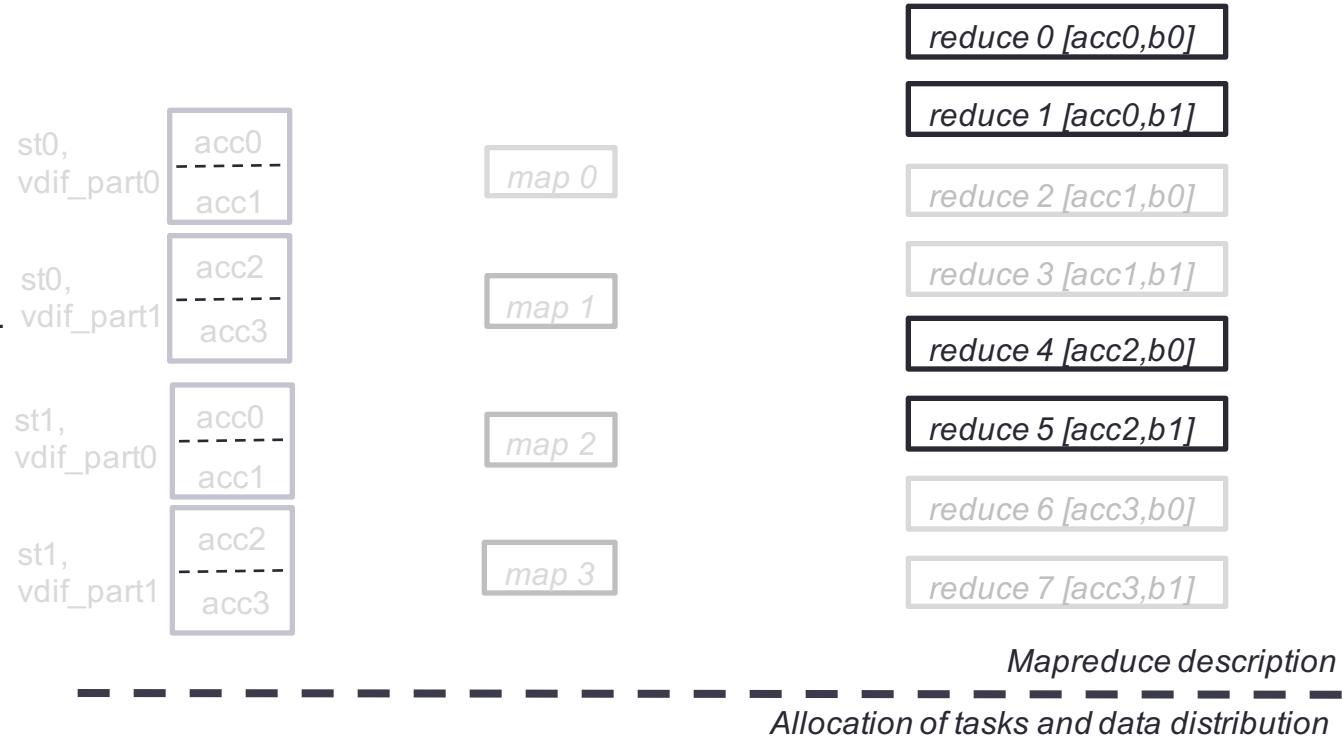
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts]*[splits]=4$
Number of reducers: $[accs]*[bands]=8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.
6. Generation of partial results.

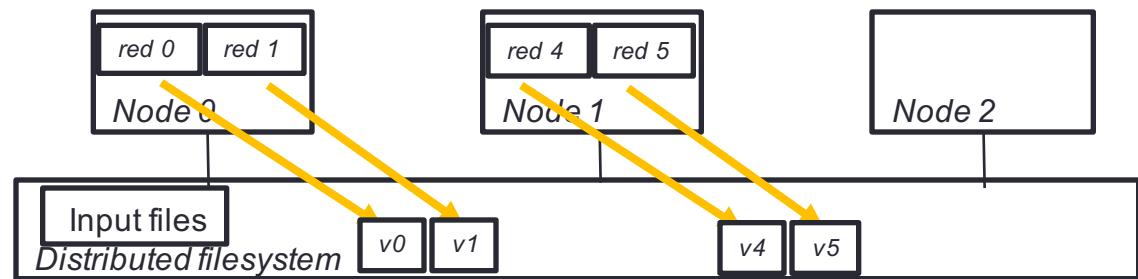


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

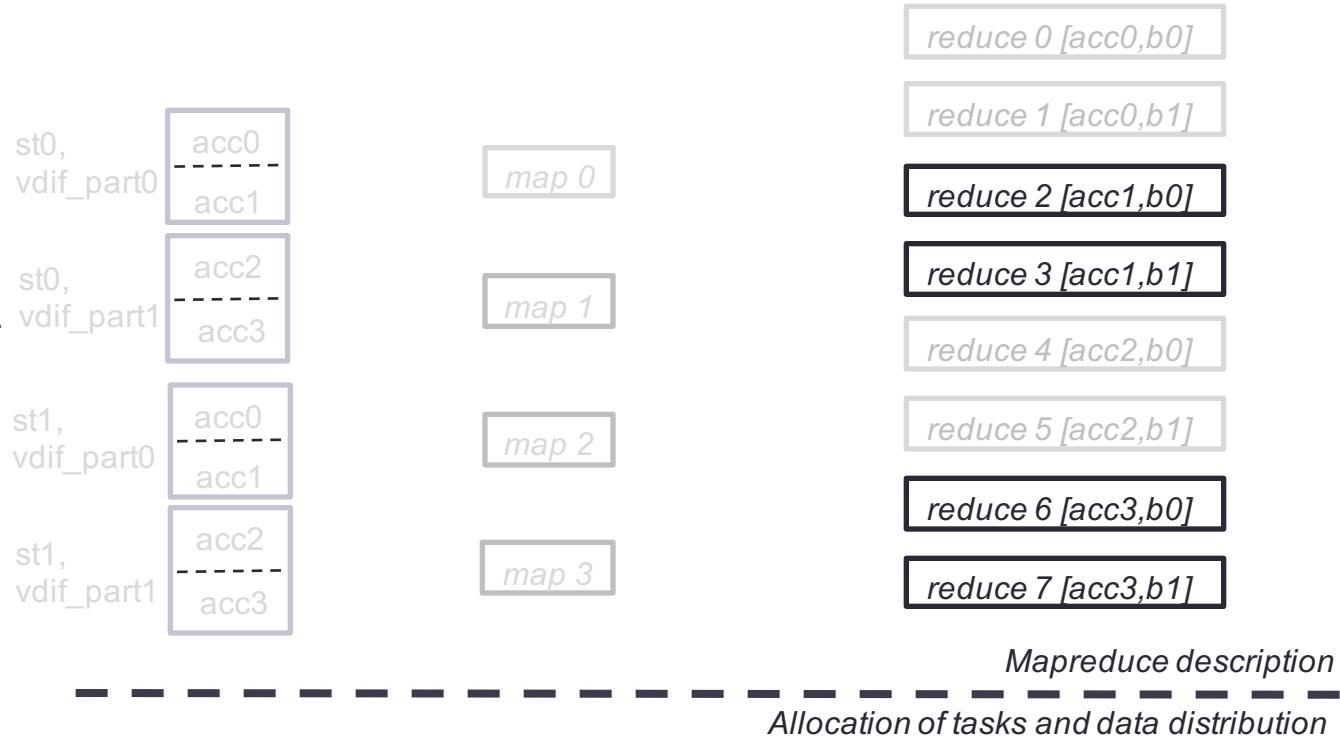
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.
6. Generation of partial results.

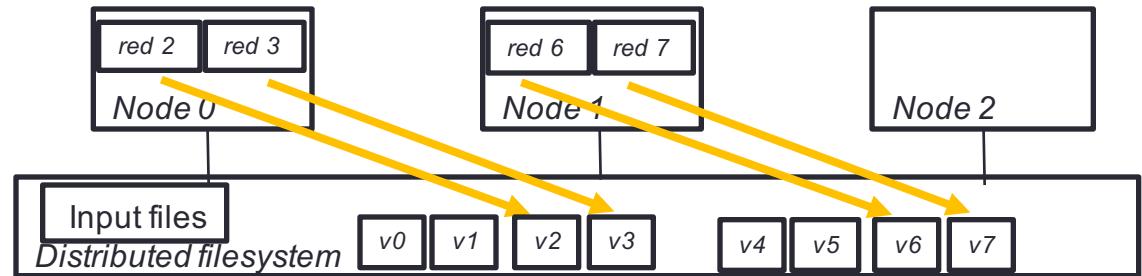


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

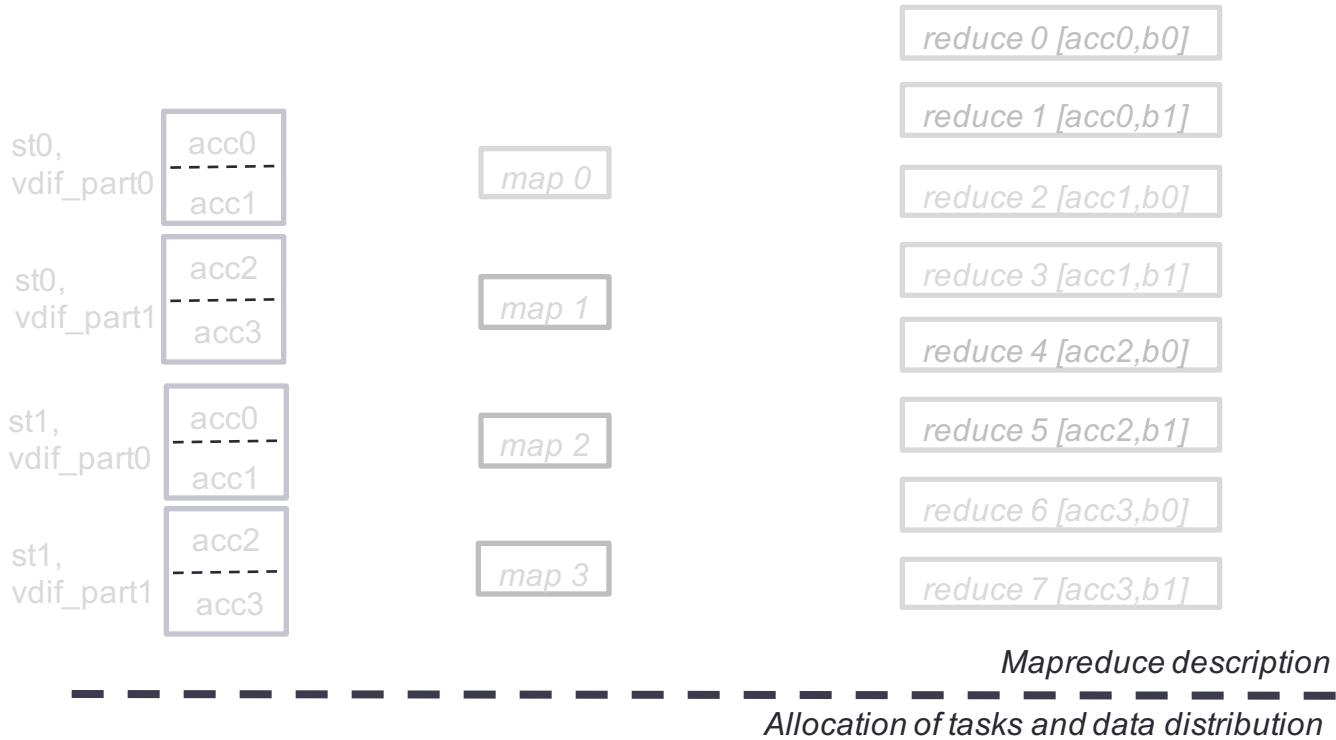
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.
6. Generation of partial results.
7. Merge results.

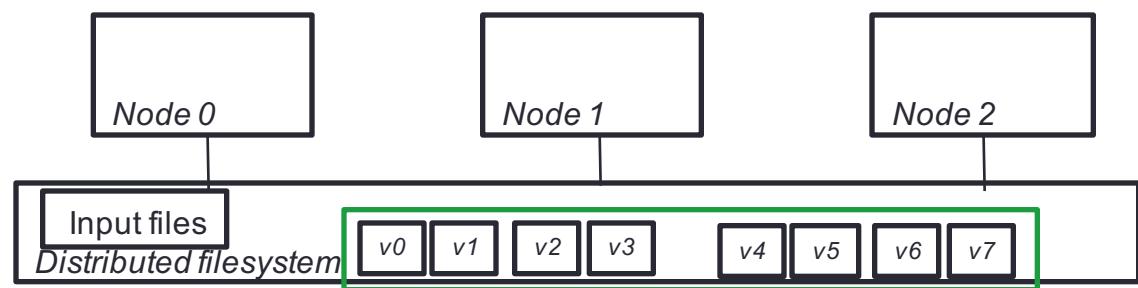


MapReduce correlation example:

Number of nodes: 3
Number of cores per node: 2

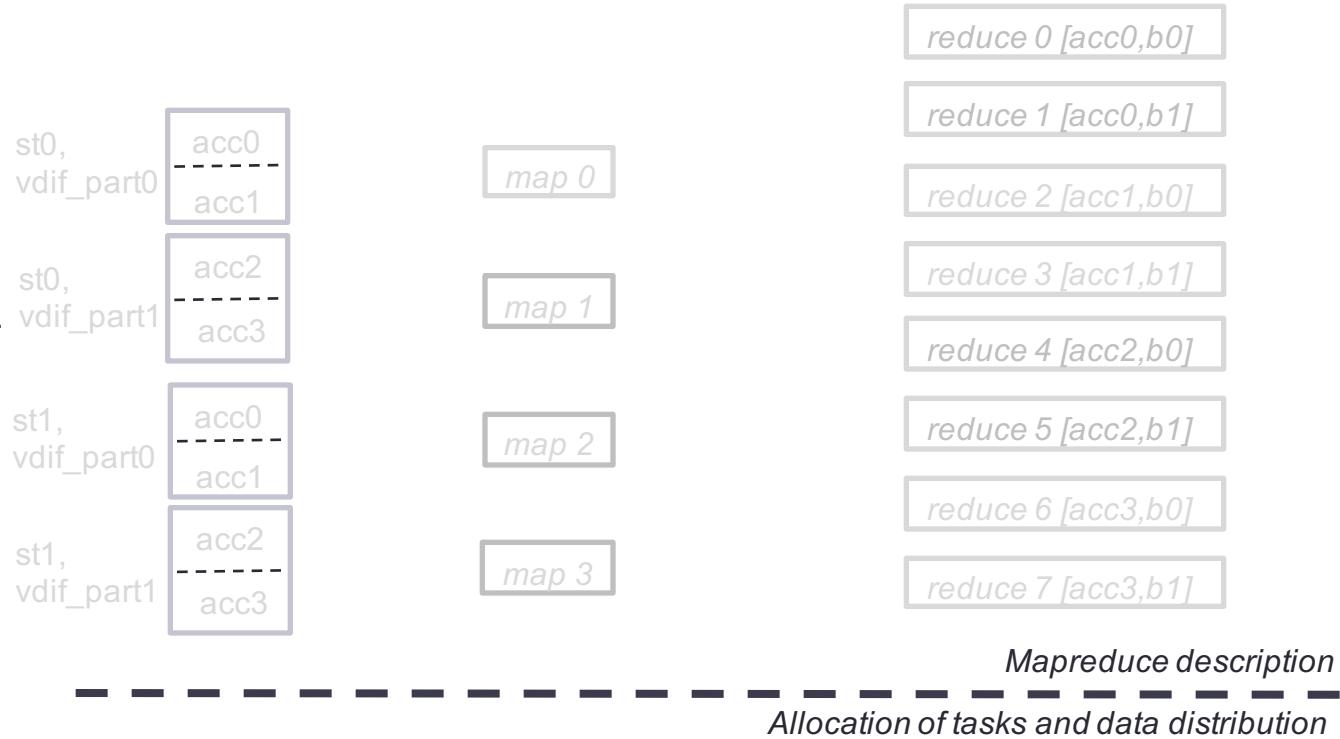
Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Correlation Example

1. Move input data into distributed filesystem.
2. The framework launches one map per inputfile.
3. Each mapper generates records with keys associated with reducers.
4. The framework distributes and sorts the records.
5. The framework will launch one reducer per partition.
6. Generation of partial results.
7. Merge results.

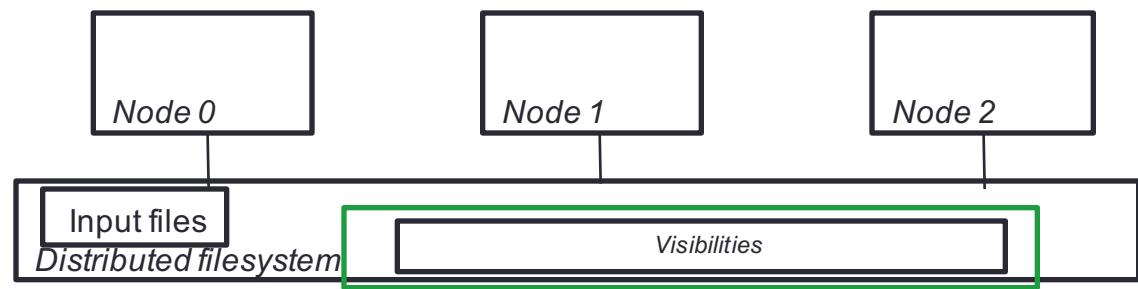


MapReduce correlation example:

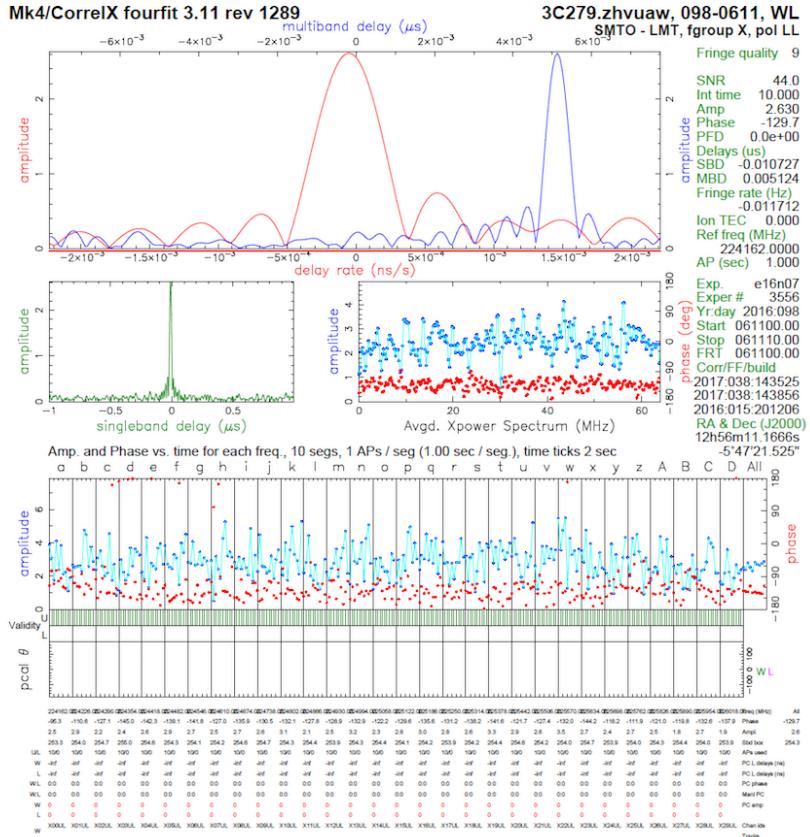
Number of nodes: 3
Number of cores per node: 2

Stations: 2
Input splits: 2 per file.
Acc. Periods: 4.
Bands: 2.

Number of mappers: $[sts] * [splits] = 4$
Number of reducers: $[accs] * [bands] = 8$



Results - alma_10s_dx_cx_20170207_1059



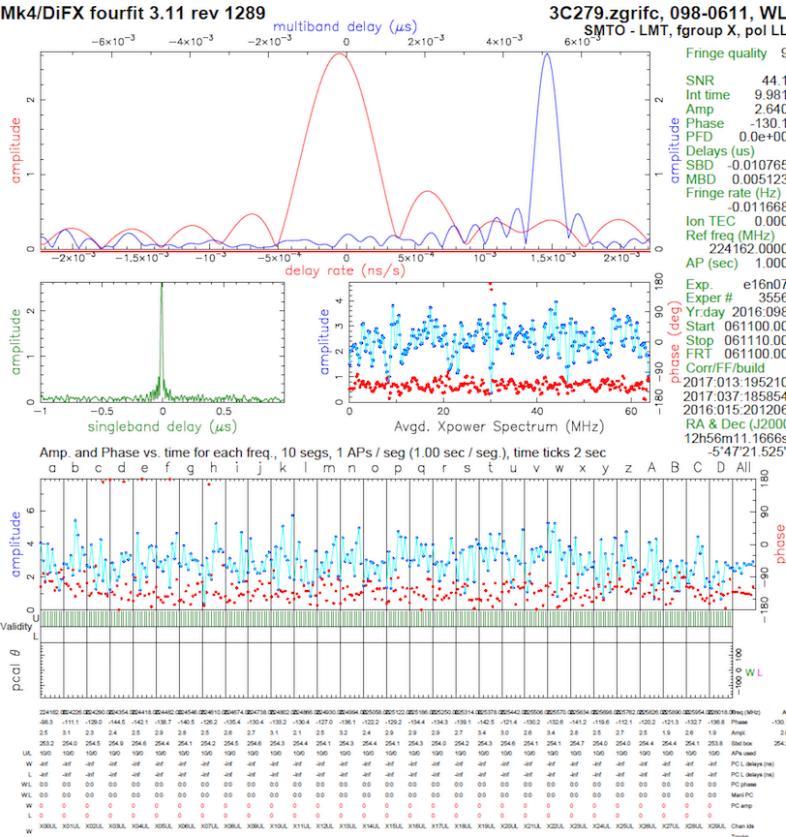
```

Group delay (usec)(model) -3.28759507128E-09 Aprior delay (usec) -3.28759506952E+03 Resid middelay (usec) 5.12423E-03 +/- 6.5E-06
Stand delay (usec) -3.28760642273E-09 Aprior clock (usec) -3.1721644E-03 Resid standdelay (usec) -0.107732E-03 +/- 2.0E-04
Stand delay (usec) -3.28759506952E+03 Aprior clockrate (/us) -1.217300E-06 Resid standdelay (usec) -0.107732E-03 +/- 3.2E-08
Delay rate (usec) 2.73221105725E-01 Aprior rate (usec) 0.2372115793E-01 Resid rate (usec) 0.234798E-09 +/- 5.6E-09
Total phase (deg) -23.71 Aprior acsol (usec/s) 1.49832132942E-05 Resid phase (deg) -129.7 +/- 2.6

RMS Theor. Amplitude 2.630 +/- 0.060 Multitone model MULTITONE, PC period (API)5, 5
ph/seg (deg) 7.4 4.1 Search (32X128) 2.554 PC rate 0.000E+00, 0.000E+00 (usec) abw window (us) -1.000 1.000
amp%/? (%) 7.0 7.2 Interp. 0.000 Bits/sample: 2x SampCntrNom: disabled mb window (us) -0.008 0.008
ph/freq (deg) 11.2 7.1 Integ. seg. avg. 2.629 Sample rate(Ms/amp): 128 dr window (ns) -0.002 0.002
ampfreq (%) 14.2 12.4 Int. freq. avg. 2.639 Data rate(MHz): 7680 nlags: 256 t_cohere infinite ion window (TC) 0.00 0.00

Lm: 15.6 E 48.7 pa:19.9 Lm: 17.2 E 65.0 pa: -6.4 u.v (fr/asec) -4091.600 4670.116 simultaneous interpolator
Control file: default Input file: /media/dl/_windows_folder_20160607/test_alpha_cx_20170210_1003/10140986-0611/W_zhuwau Output file: /media/dl/_windows_folder_20160607/test_alpha_c

```



```

Group delay (usec)(model) -3.2875057244E-03 Apron delay (usec) -3.28759569552E+03 Resid mbdelay (usec) 5.12308E-03 f, 6.5E-06
Stand delay (usec) -3.2876646027E-03 Apron clock (usec) -3.571244E-03 Resid abdelay (usec) -1.07647E-02 f, 2.0E-04
Phase delay (usec) -3.28759569552E+03 Apron microrate (use/s) 2.7322115097E-01 Resid aldelay (usec) -1.61865E-02 f, 3.2E-06
Delay rate (usec) 2.7322115097E-01 Apron rate (use/s) 2.7322115097E-01 Resid arate (usec) -5.20516E-02 f, 5.6E-09
Total phase (deg) -23.75 Apron accel (use/s) 1.49833463392E-03 Resid phase (deg) -130.1 f, -2.6

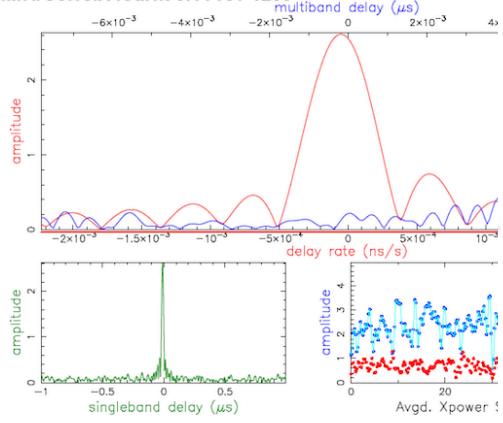
RMS Theor. Amplitude 2.54 +/- 0.060
phaseq (deg) 7.1 4.1 Search (S3X128) 5.565
phsqm(%) 8.0 7.2 Interp. 0.000
phfrq (deg) 10.7 7.1 Inc. seg. avg. 2.638
phfrq(%) 14.0 12.4 Inc. frq avg. 2.646
w_1 az: 156.0 el: 48.6 pa: -20.2 L: az: 172.6 el: 64.9 pa: -1.8 u (Hz/sec) -4078.093 465.280 simultaneous interpolator
Control file default Input file /media/f_s/windows_folder_20160907/test_alma.xml, 2017012_103925[0598-0611WL.zgrnt] Output file /media/f_s/windows_folder_20160907/test_alma.xml

```

Results - alma_10s_dx_cx_20170207_1059

Fringe quality 9

Mk4/CorrelX fourfit 3.11 rev 1289



SNR 44.0

Int time 10.000

Amp 2.630

Phase -129.7

PFD 0.0e+00

Delays (us)

SBD -0.010727

MBD 0.005124

Fringe rate (Hz)

-0.011712

Ion TEC 0.000

Ref freq (MHz)

224162.0000

AP (sec)

1.000

Exp. e16n07

Exper # 3556

Yr:day 2016:098

Start 061100.00

Stop 061110.00

FRT 061100.00

Corr/FF/build

2017:038:143525

2017:038:143856

2016:015:201206

RA & Dec (J2000)

12h56m11.1666s

-5°47'21.525"

Group delay (usec)(model) -3.28759505728E+03 Apriori delay (usec) -3.28759569552E+03
 Stand delay (usec) -3.28760642277E+03 Apriori clock (usec) -3.5712644E+00
 Phase delay (usec) -3.28759569712E+03 Apriori clockrate (usec/s) -2.7322115725E+01
 Delay rate (usec) 2.73221105725E+01 Apriori rate (usec/s) 2.7322115725E+01
 Total phase (deg) -237.1 Apriori accel (usec/s) 1.49821327492E-01

RMS Theor. Amplitude 2.630 +/- 0.060 Peal mode: MULTITON

ph/sig (deg) 7.4 4.1 Search (32X128) 2.554 Peal rate: 0.000E+00,

amp/sig (%) 7.0 7.2 Interp. 0.000 Bits/sample: 2x2

ph/rho (deg) 11.2 7.1 Inc. seg. avg. 2.629 Sample rate(MSamp/s) 7680

amp/rho (%) 14.2 12.4 Inc. sig. avg. 2.539 Data rate(Mbit/s) 76

W: az 156.3 el 48.7 pa -19.9 L: az 173.2 el 65.0 pa -6.4 u.v (frasec) .4091.600 4607.116

Control file: default Input file: /mediasrf_windows_folder_20160607/test_alma_cx_20170203_100310149896-06

Fringe quality 9

SNR 44.1

Int time 9.981

Amp 2.640

Phase -130.1

PFD 0.0e+00

Delays (us)

SBD -0.010765

MBD 0.005123

Fringe rate (Hz)

-0.011668

Ion TEC 0.000

Ref freq (MHz)

224162.0000

AP (sec)

1.000

Exp. e16n07

Exper # 3556

Yr:day 2016:098

Start 061100.00

Stop 061110.00

FRT 061100.00

Corr/FF/build

2017:013:195210

2017:037:185854

2016:015:201206

RA & Dec (J2000)

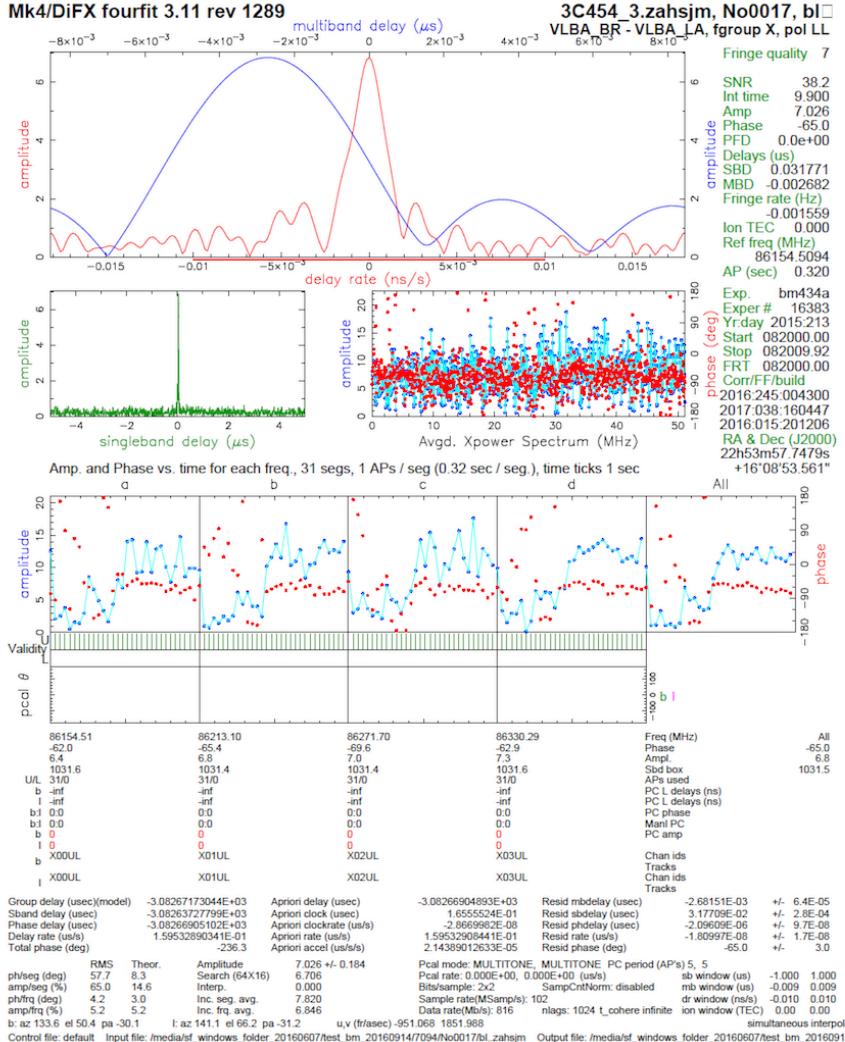
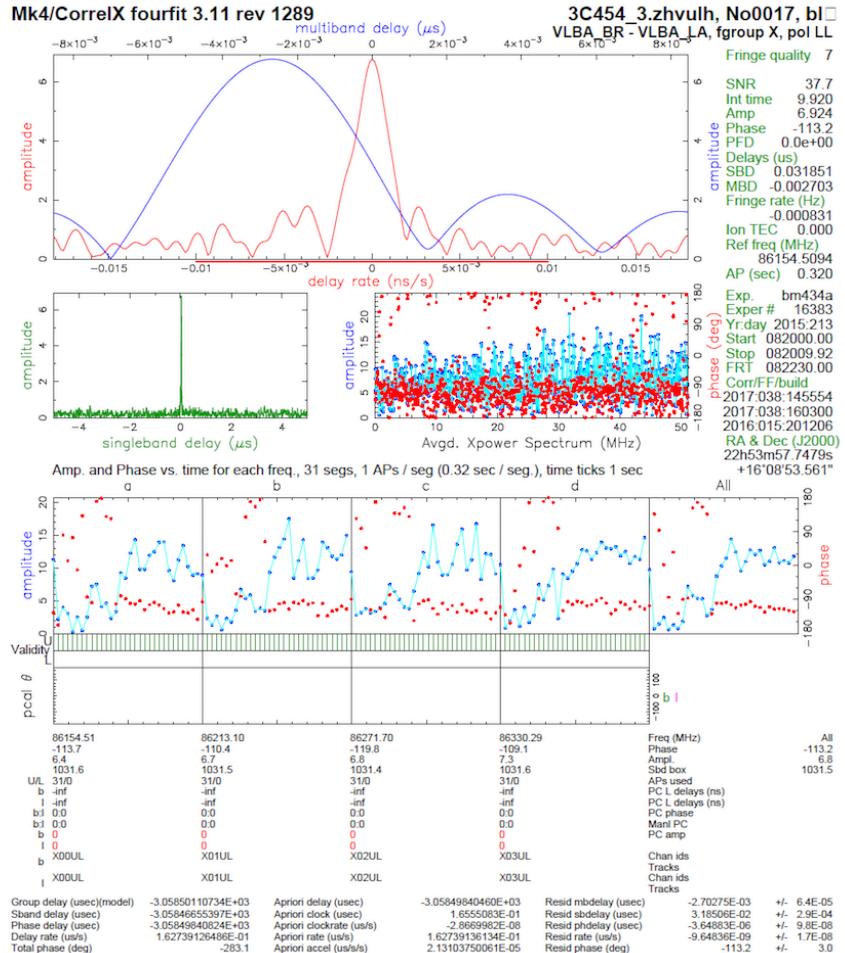
12h56m11.1666s

-5°47'21.525"

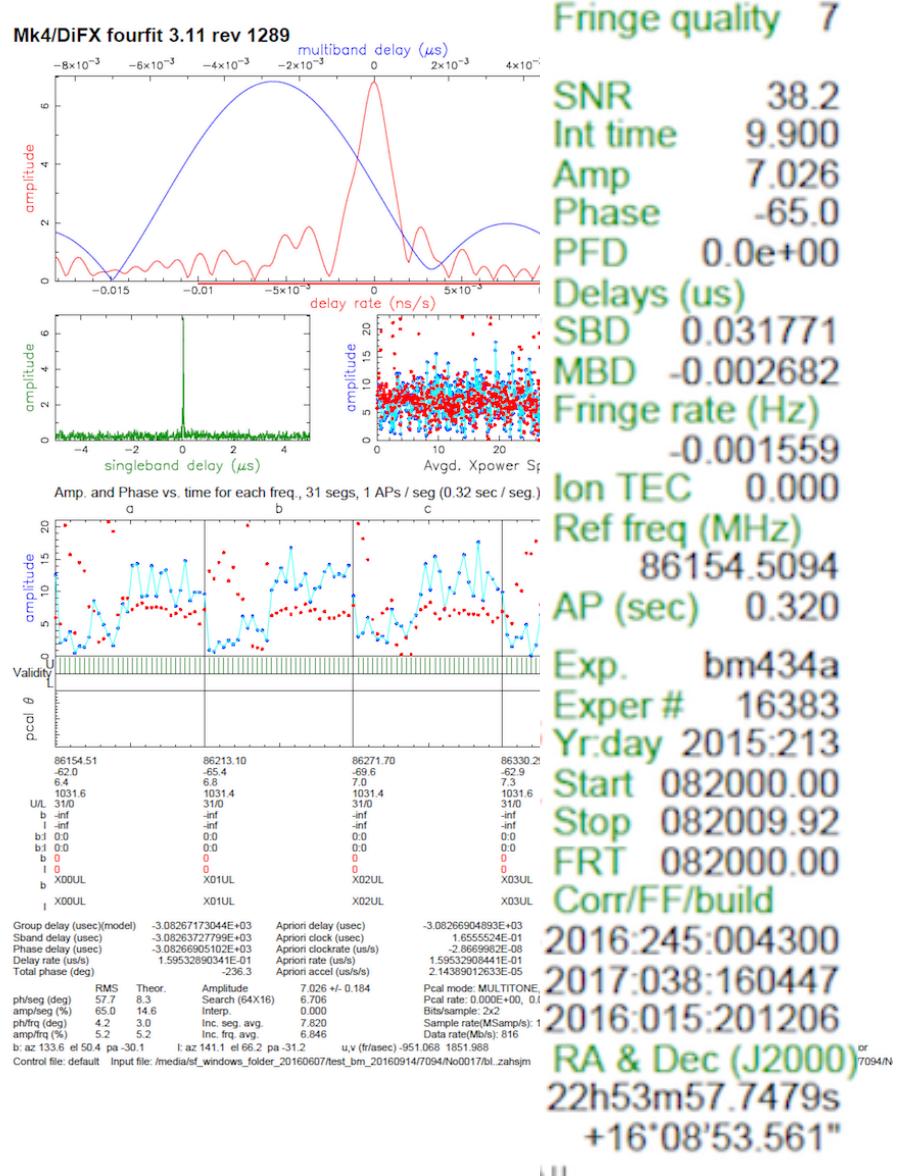
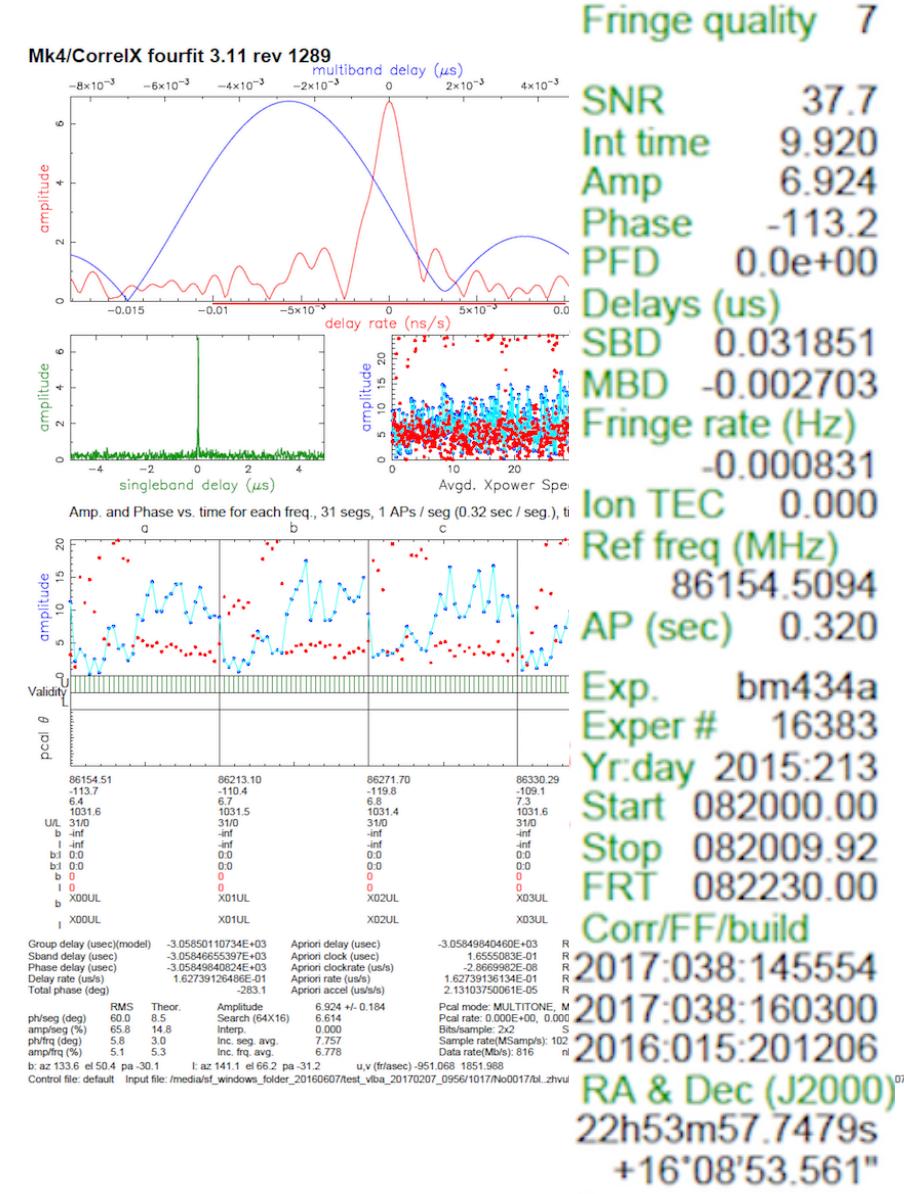


Massachusetts
Institute of
Technology

Results - vlba_10s_dx_cx_20170207_1109

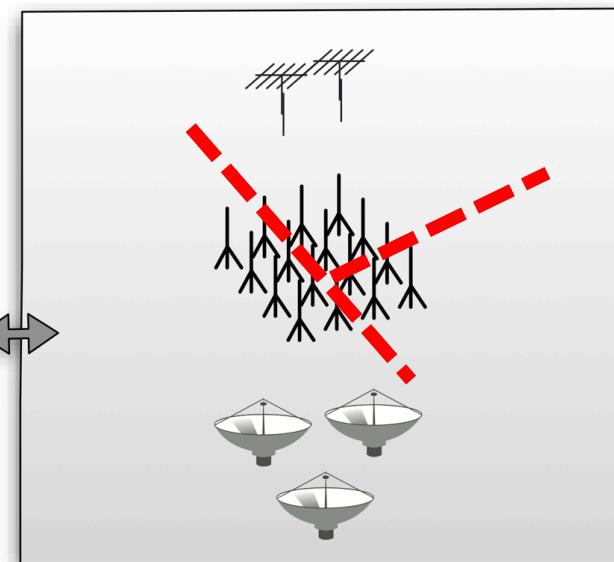
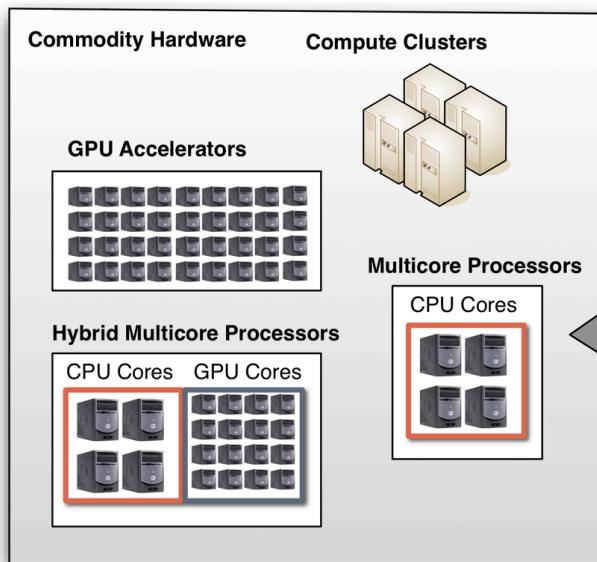


Results - vlba_10s_dx_cx_20170207_1109



Perspectives

Chunk model enables partitioning of large instrument arrays to work in different modes at the same time.



Chunk model enables execution on heterogeneous hardware. Different hardware → different reducers for chunks. Will use machine learning for performance tuning of chunk execution on each platform.

Conclusion

- CorrelX
 - **Flexibility, scalability, simplicity**
 - Proof-of-concept functionality for VGOS, ALMA, VLBA correlations
 - VGOS: LSB 2-bit complex, phase calibration
 - EHT: USB 2-bit real, zoom bands
 - Integrated in state-of-the-art processing chain
 - Plugins: Opportunity for research and teaching
- Easily prototype new algorithms, mapper-reducers, execute on real data sets

Acknowledgements

- We would like to thank the MIT Haystack Observatory staff, in particular A. **Rogers** and R. **Capallo** for invaluable help during developing and testing, and J. **Barrett**, G. **Crew**, M. **Gowanlock**, J. **Li**, G. **Rongier**, C. **Rude**, and M. **Titus** for generous feedback. We acknowledge partial support from the **NASA Space Geodesy Program** and the **NSF** for supporting the Massachusetts Green High Performance Computing Center (**MGHPC**) for our experiments in Holyoke, Massachusetts.

Thanks!

Questions? Comments? Want to contribute?

correlX@haystack.mit.edu